



Red Hat Enterprise Linux 8

Building, running, and managing containers

Building, running, and managing Linux containers on Red Hat Enterprise Linux 8

Red Hat Enterprise Linux 8 Building, running, and managing containers

Building, running, and managing Linux containers on Red Hat Enterprise Linux 8

Legal Notice

Copyright © 2020 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

Java[®] is a registered trademark of Oracle and/or its affiliates.

XFS[®] is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL[®] is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js[®] is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack[®] Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

This guide describes how to work with Linux containers on RHEL 8 systems using command-line tools such as podman, buildah, skopeo and runc.

Table of Contents

PREFACE	5
PROVIDING FEEDBACK ON RED HAT DOCUMENTATION	6
CHAPTER 1. STARTING WITH CONTAINERS	7
1.1. RUNNING CONTAINERS WITHOUT DOCKER	7
1.2. CHOOSING A RHEL ARCHITECTURE FOR CONTAINERS	8
1.3. GETTING CONTAINER TOOLS	8
1.4. RUNNING CONTAINERS AS ROOT OR ROOTLESS	9
1.4.1. Set up for rootless containers	9
1.4.2. Upgrade to rootless containers	10
1.4.3. Special considerations for rootless	11
CHAPTER 2. WORKING WITH CONTAINER IMAGES	13
2.1. SEARCHING FOR CONTAINER IMAGES	13
2.2. PULLING IMAGES FROM REGISTRIES	15
2.3. INVESTIGATING IMAGES	16
2.3.1. Listing images	16
2.3.2. Inspecting local images	16
2.3.3. Inspecting remote images	18
2.4. TAGGING IMAGES	19
2.5. SAVING AND LOADING IMAGES	19
2.6. REMOVING IMAGES	20
CHAPTER 3. WORKING WITH CONTAINERS	21
3.1. RUNNING CONTAINERS	21
3.2. INVESTIGATING RUNNING AND STOPPED CONTAINERS	23
3.2.1. Listing containers	23
3.2.2. Inspecting containers	23
3.2.3. Investigating within a container	24
3.3. STARTING AND STOPPING CONTAINERS	26
3.3.1. Starting containers	26
3.3.2. Stopping containers	26
3.4. REMOVING CONTAINERS	26
CHAPTER 4. USING RED HAT UNIVERSAL BASE IMAGES (STANDARD, MINIMAL, AND RUNTIMES)	28
4.1. WHAT ARE RED HAT BASE IMAGES?	28
4.1.1. Using standard Red Hat base images	28
4.1.2. Using minimal Red Hat base images	29
4.1.3. Using Init Red Hat base images	30
4.2. HOW ARE UBI IMAGES DIFFERENT?	30
4.3. GET UBI IMAGES	31
4.4. PULL UBI IMAGES	31
4.5. REDISTRIBUTING UBI IMAGES	31
4.6. RUN UBI IMAGES	32
4.7. ADD SOFTWARE TO A RUNNING UBI CONTAINER	32
4.7.1. Adding software to a UBI container (subscribed host)	32
4.7.2. Adding software inside the standard UBI container	32
4.7.3. Adding software inside the minimal UBI container	33
4.7.4. Adding software to a UBI container (unsubscribed host)	34
4.7.5. Build a UBI-based image	34
4.7.6. Using AppStream runtime images	36

4.7.7. Getting UBI Container Image Source Code	36
4.7.8. Tips and tricks for using UBI images	38
4.7.9. How to request new features in UBI?	38
4.7.10. How to file a support case for UBI?	38
CHAPTER 5. RUNNING SPECIAL CONTAINER IMAGES	39
5.1. TROUBLESHOOTING CONTAINER HOSTS WITH TOOLBOX	39
5.1.1. Opening privileges to the host	41
5.2. RUNNING CONTAINERS WITH RUNLABELS	41
5.2.1. Running rsyslog with runlabels	42
5.2.2. Running support-tools with runlabels	43
CHAPTER 6. BUILDING CONTAINER IMAGES WITH BUILDDAH	45
6.1. UNDERSTANDING BUILDDAH	45
6.1.1. Installing Buildah	46
6.2. GETTING IMAGES WITH BUILDDAH	46
6.3. BUILDING AN IMAGE FROM A DOCKERFILE WITH BUILDDAH	47
6.3.1. Running the image you built	47
6.3.2. Inspecting a container with Buildah	48
6.4. MODIFYING A CONTAINER TO CREATE A NEW IMAGE WITH BUILDDAH	48
6.4.1. Using buildah mount to modify a container	49
6.4.2. Using buildah copy and buildah config to modify a container	49
6.5. CREATING IMAGES FROM SCRATCH WITH BUILDDAH	50
6.6. REMOVING IMAGES OR CONTAINERS WITH BUILDDAH	51
6.7. USING CONTAINER REGISTRIES WITH BUILDDAH	52
6.7.1. Pushing containers to a private registry	52
6.7.2. Pushing containers to the Docker Hub	53
CHAPTER 7. RUNNING CONTAINERS AS SYSTEMD SERVICES WITH PODMAN	54
7.1. STARTING CONTAINERS WITH SYSTEMD	54
7.2. STARTING SERVICES WITHIN A CONTAINER USING SYSTEMD	56
CHAPTER 8. CONTAINER COMMAND-LINE REFERENCE	58
8.1. PODMAN	58
8.1.1. Using podman commands	58
8.1.2. Trying basic podman commands	60
8.1.3. Pull a container image to the local system	60
8.1.4. List local container images	60
8.1.5. Inspect a container image	60
8.1.6. Run a container image	61
8.1.7. List containers that are running or have exited	61
8.1.8. Remove a container or image	61
8.1.9. Generate a Kube pod yaml file	61
8.1.10. Generate a systemd unit file	63
8.1.11. Creating SELinux policies for containers	63
8.1.12. Using podman with MPI	63
8.1.13. Creating and restoring container checkpoints	65
8.1.13.1. Creating and restoring a container checkpoint locally	65
8.1.13.2. Reducing startup time using container restore	67
8.1.13.3. Migrating containers among systems	68
8.2. RUNC	69
8.2.1. Running containers with runc	70
8.3. SKOPEO	71
8.3.1. Inspecting container images with skopeo	71

8.3.2. Copying container images with skopeo	72
8.3.3. Getting image layers with skopeo	73
CHAPTER 9. ADDITIONAL RESOURCES	74

PREFACE

Red Hat classifies container use cases into two distinct groups: single node and multi-node, with multi-node sometimes called distributed systems. OpenShift was built to provide public, scalable deployments of containerized applications. Beyond OpenShift, however, it is useful to have a small, nimble set of tools for working with containers.

The set of container tools we are referring to can be used in a single-node use case. However, you can also wire these tools into existing build systems, CI/CD environments, and even use them to tackle workload-specific use cases, such as big data. To target the single-node use case, Red Hat Enterprise Linux (RHEL) 8 offers a set of tools to find, run, build, and share individual containers.

This guide describes how to work with Linux containers on RHEL 8 systems using command-line tools such as podman, buildah, skopeo and runc. In addition to these tools, Red Hat provides base images, to act as the foundation for your own images. Some of these base images target use cases ranging from business applications (such as Node.js, PHP, Java, and Python) to infrastructure (such as logging, data collection, and authentication).

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION

We appreciate your input on our documentation. Please let us know how we could make it better. To do so:

- For simple comments on specific passages:
 1. Make sure you are viewing the documentation in the *Multi-page HTML* format. In addition, ensure you see the **Feedback** button in the upper right corner of the document.
 2. Use your mouse cursor to highlight the part of text that you want to comment on.
 3. Click the **Add Feedback** pop-up that appears below the highlighted text.
 4. Follow the displayed instructions.
- For submitting more complex feedback, create a Bugzilla ticket:
 1. Go to the [Bugzilla](#) website.
 2. As the Component, use **Documentation**.
 3. Fill in the **Description** field with your suggestion for improvement. Include a link to the relevant part(s) of documentation.
 4. Click **Submit Bug**.

CHAPTER 1. STARTING WITH CONTAINERS

Linux Containers have emerged as a key open source application packaging and delivery technology, combining lightweight application isolation with the flexibility of image-based deployment methods.

Red Hat Enterprise Linux implements Linux Containers using core technologies such as Control Groups (Cgroups) for Resource Management, Namespaces for Process Isolation, SELinux for Security, enabling secure multi-tenancy and reducing the potential for security exploits. All this is meant to provide you with an environment to producing and running enterprise-quality containers.

Red Hat OpenShift provides powerful command-line and Web UI tools for building, managing and running containers in units referred to as **Pods**. However, there are times when you might want to build and manage individual containers and [container images](#) outside of OpenShift. Tools provided to perform those tasks that run directly on RHEL systems are described in this guide.

Unlike other container tools implementations, tools described here do not center around the monolithic Docker [container engine](#) and **docker** command. Instead, we provide a set of command-line tools that can operate without a container engine. These include:

- **podman** - For directly managing pods and container images (run, stop, start, ps, attach, exec, and so on)
- **buildah** - For building, pushing and signing container images
- **skopeo** - For copying, inspecting, deleting, and signing images
- **runc** - For providing container run and build features to podman and buildah

Because these tools are compatible with the Open Container Initiative (OCI), they can be used to manage the same Linux containers that are produced and managed by Docker and other OCI-compatible container [engines](#). However, they are especially suited to run directly on Red Hat Enterprise Linux, in single-node use cases.

For a multi-node container platform, see [OpenShift](#). Instead of relying on the single-node, daemonless tools described in this document, OpenShift requires a daemon-based container engine. Please see [Using the CRI-O Container Engine](#) for details.

1.1. RUNNING CONTAINERS WITHOUT DOCKER

Red Hat did not just remove the Docker container engine from OpenShift. It also removed the Docker container engine, along with the **docker** command, from Red Hat Enterprise Linux 8 entirely. For RHEL 8, Docker is not included and not supported by Red Hat (although it is still available from other sources).

The removal of Docker reflects a change in Red Hat's way of thinking about how containers are handled:

- In the enterprise, the focus is not on running individual containers from the command line. The primary venue for running containers is a Kubernetes-based platform, such as OpenShift.
- By repositioning OpenShift as the project for running containers, container engines like Docker become just another component of OpenShift with no direct access by end users.
- Because the container engine in OpenShift is not meant to be used directly, it can be implemented with a limited feature set that focuses on doing everything that OpenShift needs, without having to implement lots of standalone features.

Although Docker is gone from RHEL 8, and OpenShift's container engine is disconnected from single-node uses, people still want to use commands to work with containers and images manually. So Red Hat set about to create a set of tools to implement most of what the **docker** command does.

Tools like **podman**, **skopeo**, and **buildah** were developed to take over those **docker** command features. Each tool in this scenario can be more light-weight and focused on a subset of features. And with no need for a daemon process running to implement a container engine, these tools can run without the overhead of having to work with a daemon process.

If you feel that you still want to use Docker in RHEL 8, know that you can get Docker from different upstream projects, but that its use is unsupported in RHEL 8. Because so many **docker** command-line features have been implemented exactly in **podman**, you can set up an alias so that typing **docker** causes podman to run.

Installing the podman-docker package sets up such an alias. So every time you run a **docker** command line, it actually runs **podman** for you. More on this package later.

1.2. CHOOSING A RHEL ARCHITECTURE FOR CONTAINERS

Red Hat provides container images and container-related software for the following computer architectures:

- AMD64 and Intel 64 (base and layered images) (no support for the 32-bit AMD and Intel architecture)
- PowerPC 8 and 9 64-bit (base image and most layered images)
- IBM Z (base image and most layered images)
- ARM 64-bit (base image only)

Although not all Red Hat images were supported across all architectures at first, nearly all are now available on all listed architectures. See [Universal Base Images \(UBI\): Images, repositories, and packages](#) for a list of supported images.

1.3. GETTING CONTAINER TOOLS

To get an environment where you can manipulate individual containers, you can install a Red Hat Enterprise Linux 8 system, then add a set of container tools to find, run, build and share containers. Here are examples of container-related tools you can install with RHEL 8:

- **podman** - Client tool for managing containers. Can replace most features of the **docker** command for working with individual containers and images.
- **buildah** - Client tool for building OCI-compliant container images.
- **skopeo** - Client tool for copying container images to and from container registries. Includes features for signing and authenticating images as well.
- **runc** - Container runtime client for running and working with Open Container Initiative (OCI) format containers.

Using the RHEL subscription model, if you want to create container images, you must properly register and entitle the host computer on which you build them. When you install packages, as part of the process of building a container, the build process automatically has access to entitlements available from the RHEL host. So it can get RPM packages from any repository enabled on that host.

1. **Install RHEL:** If you are ready to begin, you can start by installing a Red Hat Enterprise Linux system.
2. **Register RHEL:** Once RHEL is installed, register the system. You will be prompted to enter your user name and password. Note that the user name and password are the same as your login credentials for Red Hat Customer Portal.

```
# subscription-manager register
Registering to: subscription.rhsm.redhat.com:443/subscription
Username: *****
Password: *****
```

3. **Subscribe RHEL:** Either auto subscribe or determine the pool ID of a subscription that includes Red Hat Enterprise Linux. Here is an example of auto-attaching a subscription:

```
# subscription-manager attach --auto
```

4. **Install packages:** To start building and working with individual containers, install the container-tools module, which pulls in the full set of container software packages:

```
# yum module install -y container-tools
```

5. **Install podman-docker (optional):** If you are comfortable with the **docker** command or use scripts that call **docker** directly, you can install the podman-docker package. That package installs a link that replaces the **docker** command-line interface with the matching **podman** commands instead. It also links the man pages together, so **man docker info** will show the **podman info** man page.

```
# yum install -y podman-docker
```

1.4. RUNNING CONTAINERS AS ROOT OR ROOTLESS

Running the container tools such as **podman**, **skopeo**, or **buildah** as a user with superuser privilege (root user) is the best way to ensure that your containers have full access to any feature available on your system. However, with the feature called "Rootless Containers," generally available as of RHEL 8.1, you can work with containers as a regular user.

Although container engines, such as Docker, let you run **docker** commands as a regular (non-root) user, the docker daemon that carries out those requests runs as root. So, effectively, regular users can make requests through their containers that harm the system, without there being clarity about who made those requests. By setting up rootless container users, system administrators limit potentially damaging container activities from regular users, while still allowing those users to safely run many container features under their own accounts.

This section describes how to set up your system to use container tools (Podman, Skopeo, and Buildah) to work with containers as a non-root user (rootless). It also describes some of the limitations you will encounter because regular user accounts don't have full access to all operating system features that their containers might need to run.

1.4.1. Set up for rootless containers

You need to become root user to set up your RHEL system to allow non-root user accounts to use container tools:

1. **Install RHEL:** Install RHEL 8.1 or upgrade to RHEL 8.1 from RHEL 8.0. Earlier RHEL 7 versions are missing features needed for this procedure. If you are upgrading from RHEL 7.6 or earlier, continue to "Upgrade to rootless containers" after this procedure is done.
2. **Install podman and slirp4netns:** If not already installed, install the podman and slirp4netns packages:

```
# yum install slirp4netns podman -y
```

3. **Increase user namespaces:** To increase the number of user namespaces in the kernel, type the following:

```
# echo "user.max_user_namespaces=28633" > /etc/sysctl.d/usersns.conf
# sysctl -p /etc/sysctl.d/usersns.conf
```

4. **Create a new user account:** To create a new user account and add a password for that account (for example, joe), type the following:

```
# useradd -c "Joe Jones" joe
# passwd joe
```

The user is automatically configured to be able to use rootless podman.

5. **Try a podman command:** Log in directly as the user you just configured (don't use **su** or **su -** to become that user because that doesn't set the correct environment variables) and try to pull and run an image:

```
$ podman pull registry.access.redhat.com/ubi8/ubi
$ podman run registry.access.redhat.com/ubi8/ubi cat /etc/os-release
NAME="Red Hat Enterprise Linux"
VERSION="8.1 (Ootpa)"
...
```

6. **Check rootless configuration:** To check that your rootless configuration is set up properly, you can run commands inside the modified user namespace with the **podman unshare** command. As the rootless user, the following command lets you see how the uids are assigned to the user namespace:

```
$ podman unshare cat /proc/self/uid_map
0    1001    1
1    65537   65536
```

1.4.2. Upgrade to rootless containers

If you have upgraded from RHEL 7, you must configure subuid and subgid values manually for any existing user you want to be able to use rootless podman.

Using an existing user name and group name (for example, jill), set the range of accessible user and group IDs that can be used for their containers. Here are a couple of warnings:

- Don't include the rootless user's UID and GID in these ranges
- If you set multiple rootless container users, use unique ranges for each user

- We recommend 65536 UIDs and GIDs for maximum compatibility with existing container images, but the number can be reduced
- Never use UIDs or GIDs under 1000 or reuse UIDs or GIDs from existing user accounts (which, by default, start at 1000)
Here is an example:

```
# echo "jill:165537:65536" >> /etc/subuid
# echo "jill:165537:65536" >> /etc/subgid
```

The user/group jill is now allocated 65535 user and group IDs, ranging from 165537-231072. That user should be able to begin running commands to work with containers now.

1.4.3. Special considerations for rootless

Here are some things to consider when running containers as a non-root user:

- As a non-root container user, container images are stored under your home directory (**\$HOME/.local/share/containers/storage/**), instead of **/var/lib/containers**.
- Users running rootless containers are given special permission to run as a range of user and group IDs on the host system. However, they otherwise have no root privileges to the operating system on the host.
- If you need to configure your rootless container environment, edit configuration files in your home directory (**\$HOME/.config/containers**). Configuration files include **storage.conf** (for configuring storage) and **libpod.conf** (for a variety of container settings). You could also create a **registries.conf** file to identify container registries available when you use **podman** to pull, search or run images.
- A container running as root in a rootless account can turn on privileged features within its own namespace. But that doesn't provide any special privileges to access protected features on the host (beyond having extra UIDs and GIDs). Here are examples of container actions you might expect to work from a rootless account that will not work:
 - Anything you want to access from a mounted directory from the host must be accessible by the UID running your container or your request to access that component will fail.
 - There are some system features you won't be able to change without privilege. For example, you cannot change the system clock by simply setting a SYS_TIME capability inside a container and running the network time service (ntpd). You would have to run that container as root, bypassing your rootless container environment and using the root user's environment, for that capability to work, such as:

```
$ sudo podman run -d --cap-add SYS_TIME ntpd
```

Note that this example allows ntpd to adjust time for the entire system, and not just within the container.

- A rootless container has no ability to access a port less than 1024. Inside the rootless container's namespace it can, for example, start a service that exposes port 80 from an httpd service from the container, but it will not be accessible outside of the namespace:

```
$ podman run -d httpd
```

However, a container would need root privilege, again using the root user's container environment, to expose that port to the host system:

```
$ sudo podman run -d -p 80:80 httpd
```

- The administrator of a workstation can configure it to allow users to expose services below 1024, but they should understand the security implications. A regular user could, for example, run a web server on the official port 80 and trick external users into believing that it was configured by the administrator. This is generally OK on a workstation, but might not be on a network-accessible development server, and definitely should not be done on production servers. To allow users to bind to ports down to port 80 run the following command:

```
# echo 80 > /proc/sys/net/ipv4/ip_unprivileged_port_start
```

- Rootless containers currently relies on setting static subuid and subgid ranges. If you are using LDAP or Active Directory to provide user authentication, there is no automated way to provide those UID and GID ranges to users. A current workaround could be to set static ranges in `/etc/subuid` and `/etc/subgid` files to match the known UIDs and GIDs in use.
- Container storage must be on a local file system, because remote file systems do not work well with unprivileged user namespaces.
- An on-going list of shortcomings of running **podman** and related tools without root privilege is contained in [Shortcomings of Rootless Podman](#).

CHAPTER 2. WORKING WITH CONTAINER IMAGES

Using **podman**, you can run, investigate, start, stop, investigate, and remove container images.

2.1. SEARCHING FOR CONTAINER IMAGES

The **podman search** command lets you search selected container registries for images.



NOTE

You can also search for images in the [Red Hat Container Registry](#). The Red Hat Container Registry includes the image description, contents, health index, and other information.

You can find the list of registries in the configuration file **registries.conf**:

```
[registries.search]
registries = ['registry.access.redhat.com', 'registry.redhat.io', 'docker.io']

[registries.insecure]
registries = []

[registries.block]
registries = []
```

- By default, the **podman search** command searches for container images from registries listed in section **[registries.search]** in the given order. In this case, **podman search** command looks for the requested image in registry.access.redhat.com, registry.redhat.io and docker.io in this order.
- The **[registries.insecure]** section adds the registries that do not use TLS (an insecure registry).
- The **[registries.block]** section disallows the access to the registry from your local system.

As a root user, you can edit the **/etc/containers/registries.conf** file to change the default, system-wide search settings.

As a regular (rootless) user of **podman**, you can create your own **registries.conf** file in your home directory (**\$HOME/.config/containers/registries.conf**) to override the system-wide settings.

Make sure that you follow the conditions when configuring container registries:

- Each registry must be surrounded by single quotes.
- If there are multiple registries set for the **registries = value**, you must separate those registries by commas.
- You can identify registries by either IP address or hostname.
- If the registry uses a non-standard port – other than TCP ports 443 for secure and 80 for insecure, enter that port number with the registry name. For example: host.example.com:9999.
- The system searches for registries in the order in which they appear in the **registries.search** list of the **registries.conf** file.

Some **podman search** command examples follow. The first example illustrates the unsuccessful search of all images from quay.io. The forwardslash at the end means to search the whole registry for all images accessible to you:

```
# podman search quay.io/
ERRO[0000] error searching registry "quay.io": couldn't search registry "quay.io":
unable to retrieve auth token: invalid username/password
```

To search quay.io registry, log in first:

```
# podman login quay.io
Username: johndoe
Password: *****
Login Succeeded!
# podman search quay.io/
```

INDEX	NAME	DESCRIPTION	STARS	OFFICIAL	AUTOMATED
quay.io	quay.io/test/myquay		0		
quay.io	quay.io/test/redistest		0		
quay.io	quay.io/johndoe/websrv21		0		
quay.io	quay.io/johndoe/mydbtest		0		
quay.io	quay.io/johndoe/newbuild-10		0		

Search all available registries for **postgresql** images (resulting in more than 40 images found):

```
# podman search postgresql-10
```

INDEX	NAME	DESCRIPTION	STARS	OFFICIAL	AUTOMATED
redhat.io	registry.redhat.io/rhel8/postgresql-10	This container image ...	0		
redhat.io	registry.redhat.io/rhsc1/postgresql-10-rhel7	PostgreSQL is an advanced ...	0		
quay.io	quay.io/mettle/postgresql-database-provisioning				
docker.io	docker.io/centos/postgresql-10-centos7	PostgreSQL is an advanced ...	13		
...					

To limit your search for **postgresql** to images from registry.redhat.io, type the following command. Notice that by entering the registry and the image name, any repository in the registry can be matched:

```
# podman search registry.redhat.io/postgresql-10
```

INDEX	NAME	DESCRIPTION	STARS	OFFICIAL	AUTOMATED
redhat.io	registry.redhat.io/rhel8/postgresql-10	This container image ...	0		
redhat.io	registry.redhat.io/rhsc1/postgresql-10-rhel7	PostgreSQL is an ...	0		

To get longer descriptions for each container image, add **--no-trunc** to the command:

```
# podman search --no-trunc registry.redhat.io/rhel8/postgresql-10
```

INDEX	NAME	DESCRIPTION	STARS	OFFICIAL	AUTOMATED
redhat.io	registry.redhat.io/rhel8/postgresql-10	This container image provides a containerized packaging of the PostgreSQL postgres daemon and client application. The postgres server daemon accepts connections from clients and provides access to content from PostgreSQL databases on behalf of the clients.	0		

To access insecure registries, add the fully-qualified name of the registry to the **[registries.insecure]** section of the **/etc/containers/registries.conf** file. For example:

```
[registries.search]
registries = ['myregistry.example.com']

[registries.insecure]
registries = ['myregistry.example.com']
```

Then, search for **myimage** images:

```
# podman search myregistry.example.com/myimage
INDEX    NAME DESCRIPTION                               STARS OFFICIAL AUTOMATED
example.com myregistry.example.com/myimage
The myimage container executes the ... 0
```

Now you can pull **myimage** image:

```
# podman pull myimage.example.com/myimage
```

2.2. PULLING IMAGES FROM REGISTRIES

To get container images from a remote registry (such as Red Hat's own container registry) and add them to your local system, use the **podman pull** command:

```
# podman pull <registry>[:<port>]/[<namespace>]/<name>:<tag>
```

The *<registry>* is a host that provides a container registry service on TCP *<port>*. Together, *<namespace>* and *<name>* identify a particular image controlled by *<namespace>* at that registry. The *<tag>* is an additional name to locally-stored image, the default tag is **latest**. Always use fully qualified image names including: registry, namespace, image name and tag. When using short names, there is always an inherent risk of spoofing. Add registries that are trusted, that is registries which do not allow unknown or anonymous users to create accounts with arbitrary names.

Some registries also support raw *<name>*; for those, *<namespace>* is optional. When it is included, however, the additional level of hierarchy that *<namespace>* provides is useful to distinguish between images with the same *<name>*. For example:

Namespace	Examples (<i><namespace>/<name></i>)
organization	redhat/kubernetes, google/kubernetes
login (user name)	alice/application, bob/application
role	devel/database, test/database, prod/database

The registries that Red Hat provides are registry.redhat.io (requiring authentication), registry.access.redhat.com (requires no authentication), and registry.connect.redhat.com (holds [Red Hat Partner Connect](#) program images). For details on the transition to registry.redhat.io, see [Red Hat Container Registry Authentication](#). Before you can pull containers from registry.redhat.io, you need to authenticate. For example:

```
# podman login registry.redhat.io
Username: myusername
Password: *****
Login Succeeded!
```

Use the pull option to pull an image from a remote registry. To pull the rhel base image and rsyslog logging image from the Red Hat registry, type:

```
# podman pull registry.redhat.io/ubi8/ubi
# podman pull registry.redhat.io/rhel8/rsyslog
```

An image is identified by a registry name (registry.redhat.io), a namespace name (ubi8) and the image name (ubi). You could also add a tag (which defaults to :latest if not entered). The repository name **ubi**, when passed to the **podman pull** command without the name of a registry preceding it, is ambiguous and could result in the retrieval of an image that originates from an untrusted registry. If there are multiple versions of the same image, adding a tag, such as **latest** to form a name such as **ubi8/ubi:latest**, lets you choose the image more explicitly.

To see the images that resulted from the above **podman pull** command, along with any other images on your system, type **podman images**:

```
REPOSITORY          TAG   IMAGE ID   CREATED   SIZE
registry.redhat.io/ubi8/ubi   latest eb205f07ce7d 2 weeks ago 214MB
registry.redhat.io/rhel8/rsyslog latest 85cfba5cd49c 2 weeks ago 234MB
```

The **ubi** and **rsyslog** images are now available on your local system for you to work with.

2.3. INVESTIGATING IMAGES

Using **podman images** you can see which images have been pulled to your local system. To look at the metadata associated with an image, use **podman inspect**.

2.3.1. Listing images

To see which images have been pulled to your local system and are available to use, type:

```
# podman images
REPOSITORY          TAG   IMAGE ID   CREATED   VIRTUAL SIZE
registry.redhat.io/rhel8/support-tools latest b3d6ce4e0043 2 days ago 234MB
registry.redhat.io/ubi8/ubi-init   latest 779a05997856 2 days ago 225MB
registry.redhat.io/ubi8/ubi       latest a80dad1c1953 3 days ago 210MB
```

2.3.2. Inspecting local images

After you pull an image to your local system and before you run it, it is a good idea to investigate that image. Reasons for investigating an image before you run it include:

- Understanding what the image does
- Checking what software is inside the image

The **podman inspect** command displays basic information about what an image does. You also have the option of mounting the image to your host system and using tools from the host to investigate what's in the image. Here is an example of investigating what a container image does before you run it:

1. **Inspect an image** Run **podman inspect** to see what command is executed when you run the container image, as well as other information. Here are examples of examining the `ubi8/ubi` and `rhel8/rsyslog` container images (with only snippets of information shown here):

```
# podman pull registry.redhat.io/ubi8/ubi
# podman inspect registry.redhat.io/ubi8/ubi | less
```

```
...
  "Cmd": [
    "/bin/bash"
  ],
  "Labels": {
    "License": "GPLv3",
    "architecture": "x86_64",
    "authoritative-source-url": "registry.redhat.io",
    "build-date": "2018-10-24T16:46:08.916139",
    "com.redhat.build-host": "cpt-0009.osbs.prod.upshift.rdu2.redhat.com",
    "com.redhat.component": "rhel-server-container",
    "description": "The Red Hat Enterprise Linux Base image is designed to be a fully
supported..."
  }
  ...
```

```
# podman pull registry.redhat.io/rhel8/rsyslog
# podman inspect registry.redhat.io/rhel8/rsyslog
  "Cmd": [
    "/bin/rsyslog.sh"
  ],
  "Labels": {
    "License": "GPLv3",
    "architecture": "x86_64",
    ...
    "install": "podman run --rm --privileged -v /:/host -e HOST=/host -e IMAGE=IMAGE -e
NAME=NAME IMAGE /bin/install.sh",
    ...
    "run": "podman run -d --privileged --name NAME --net=host --pid=host -v
/etc/pki/rsyslog:/etc/pki/rsyslog -v /etc/rsyslog.conf:/etc/rsyslog.conf -v
/etc/sysconfig/rsyslog:/etc/sysconfig/rsyslog -v /etc/rsyslog.d:/etc/rsyslog.d -v /var/log:/var/log
-v /var/lib/rsyslog:/var/lib/rsyslog -v /run:/run -v /etc/machine-id:/etc/machine-id -v
/etc/localtime:/etc/localtime -e IMAGE=IMAGE -e NAME=NAME --restart=always IMAGE
/bin/rsyslog.sh",
    "summary": "A containerized version of the rsyslog utility"
  }
  ...
```

The `ubi8/ubi` container will execute the bash shell, if no other argument is given when you start it with **podman run**. If an Entrypoint were set, its value would be used instead of the `Cmd` value (and the value of `Cmd` would be used as an argument to the Entrypoint command).

In the second example, the `rhel8/rsyslog` container image has built-in **install** and **run** labels. Those labels give an indication of how the container is meant to be set up on the system (install) and executed (run).

2. **Mount a container:** Using the **podman** command, mount an active container to further investigate its contents. This example runs and lists a running **rsyslog** container, then displays the mount point from which you can examine the contents of its file system:

```
# podman run -d registry.redhat.io/rhel8/rsyslog
# podman ps
CONTAINER ID IMAGE          COMMAND             CREATED   STATUS    PORTS
NAMES
1cc92aea398d ...rsyslog:latest /bin/rsyslog.sh 37 minutes ago Up 1 day ago    myrsyslog
# podman mount 1cc92aea398d
/var/lib/containers/storage/overlay/65881e78.../merged
# ls /var/lib/containers/storage/overlay/65881e78*/merged
bin boot dev etc home lib lib64 media mnt opt proc root run sbin srv sys tmp usr
var
```

After the **podman mount**, the contents of the container are accessible from the listed directory on the host. Use **ls** to explore the contents of the image.

3. **Check the image's package list** To check the packages installed in the container, tell the **rpm** command to examine the packages installed on the container's mount point:

```
# rpm -qa --root=/var/lib/containers/storage/overlay/65881e78.../merged
redhat-release-server-7.6-4.el7.x86_64
filesystem-3.2-25.el7.x86_64
basesystem-10.0-7.el7.noarch
ncurses-base-5.9-14.20130511.el7_4.noarch
glibc-common-2.17-260.el7.x86_64
nspr-4.19.0-1.el7_5.x86_64
libstdc++-4.8.5-36.el7.x86_64
```

2.3.3. Inspecting remote images

To inspect a container image before you pull it to your system, you can use the **skopeo inspect** command. With **skopeo inspect**, you can display information about an image that resides in a remote container registry.

The following command inspects the **ubi8-init** image from the Red Hat registry:

```
# skopeo inspect docker://registry.redhat.io/ubi8/ubi8-init
{
  "Name": "registry.redhat.io/ubi8/ubi8-init",
  "Digest": "sha256:53dfe24...",
  "RepoTags": [
    "8.0.0-9",
    "8.0.0",
    "latest"
  ],
  "Created": "2019-05-13T20:50:11.437931Z",
  "DockerVersion": "1.13.1",
  "Labels": {
    "architecture": "x86_64",
    "authoritative-source-url": "registry.access.redhat.com",
    "build-date": "2019-05-13T20:49:44.207967",
```

```
"com.redhat.build-host": "cpt-0013.osbs.prod.upshift.rdu2.redhat.com",
"com.redhat.component": "ubi8-init-container",
"description": "The Red Hat Enterprise Linux Init image is designed to be..."
```

2.4. TAGGING IMAGES

You can add names to images to make it more intuitive to understand what they contain. Tagging images can also be used to identify the target registry for which the image is intended. Using the **podman tag** command, you essentially add an alias to the image that can consist of several parts. Those parts can include:

```
registryhost/username/NAME:tag
```

You can add just *NAME* if you like. For example:

```
# podman tag 474ff279782b myrhel8
```

In the previous example, the **rhel8** image had an image ID of 474ff279782b. Using **podman tag**, the name **myrhel8** now also is attached to the image ID. So you could run this container by name (rhel8 or myrhel8) or by image ID. Notice that without adding a :tag to the name, it was assigned :latest as the tag. You could have set the tag to 8.0 as follows:

```
# podman tag 474ff279782b myrhel8:8.0
```

To the beginning of the name, you can optionally add a user name and/or a registry name. The user name is actually the repository on Docker.io that relates to the user account that owns the repository. Tagging an image with a registry name was shown in the "Tagging Images" section earlier in this document. Here's an example of adding a user name:

```
# podman tag 474ff279782b jsmith/myrhel8
# podman images | grep 474ff279782b
rhel8      latest 474ff279782b 7 days ago 139.6 MB
myrhel8    latest 474ff279782b 7 months ago 139.6 MB
myrhel8    7.1    474ff279782b 7 months ago 139.6 MB
jsmith/myrhel8 latest 474ff279782b 7 months ago 139.6 MB
```

Above, you can see all the image names assigned to the single image ID.

2.5. SAVING AND LOADING IMAGES

If you want to save a container image you have stored locally, you can use **podman save** to save the image to an archive file or directory and restore it later to another container environment. The archive you save can be in any of several different container image formats: docker-archive, oci-archive, oci-dir (directory with oci manifest type), or docker-dir (directory with v2s2 manifest type). After you save an image, you can store it or send it to someone else, then **load** the image later to reuse it. Here is an example of saving an image as a tarball in the default docker-archive format:

```
# podman save -o myrsyslog.tar registry.redhat.io/rhel8/rsyslog:latest
# file myrsyslog.tar
myrsyslog.tar: POSIX tar archive
```

The **myrsyslog.tar** file is now stored in your current directory. Later, when you are ready to reuse the tarball as a container image, you can import it to another podman environment as follows:

```
# podman load -i myrsyslog.tar
# podman images
REPOSITORY          TAG   IMAGE ID   CREATED   SIZE
registry.redhat.io/rhel8/rsyslog latest 1f5313131bf0 7 weeks ago 235 MB
```

Instead of using **save** and **load** to store and reload an image, you can make a copy of a container instead, using **podman export** and **podman import**.

2.6. REMOVING IMAGES

To see a list of images that are on your system, run the **podman images** command. To remove images you no longer need, use the **podman rmi** command, with the image ID or name as an option. (You must stop any containers run from an image before you can remove the image.) Here is an example:

```
# podman rmi ubi8-init
7e85c34f126351ccb9d24e492488ba7e49820be08fe53bee02301226f2773293
```

You can remove multiple images on the same command line:

```
# podman rmi registry.redhat.io/rhel8/rsyslog support-tools
46da8e23fa1461b658f9276191b4f473f366759a6c840805ed0c9ff694aa7c2f
85cfba5cd49c84786c773a9f66b8d6fca04582d5d7b921a308f04bb8ec071205
```

If you want to clear out all your images, you could use a command like the following to remove all images from your local registry (make sure you mean it before you do this!):

```
# podman rmi -a
1ca061b47bd70141d11dcb2272dee0f9ea3f76e9afd71cd121a000f3f5423731
ed904b8f2d5c1b5502dea190977e066b4f76776b98f6d5aa1e389256d5212993
83508706ef1b603e511b1b19afcb5faab565053559942db5d00415fb1ee21e96
```

To remove images that have multiple names (tags) associated with them, you need to add the force option to remove them. For example:

```
# podman rmi -a
A container associated with containers/storage, i.e. via Buildah, CRI-O, etc., may be associated with
this image: 1de7d7b3f531

# podman rmi -f 1de7d7b3f531
1de7d7b3f531...
```


CHAPTER 3. WORKING WITH CONTAINERS

Containers represent a running or stopped process spawned from the files located in a decompressed container image. Tools for running containers and working with them are described in this section.

3.1. RUNNING CONTAINERS

When you execute a **podman run** command, you essentially spin up and create a new container from a container image. The command you pass on the **podman run** command line sees the inside of the container as its running environment so, by default, very little can be seen of the host system. For example, by default, the running application sees:

- The file system provided by the container image.
- A new process table from inside the container (no processes from the host can be seen).

If you want to make a directory from the host available to the container, map network ports from the container to the host, limit the amount of memory the container can use, or expand the CPU shares available to the container, you can do those things from the **podman run** command line. Here are some examples of **podman run** command lines that enable different features.

EXAMPLE #1 (Run a quick command): This podman command runs the **cat /etc/os-release** command to see the type of operating system used as the basis for the container. After the container runs the command, the container exits and is deleted (**--rm**).

```
# podman run --rm registry.redhat.io/ubi8/ubi cat /etc/os-release
NAME="Red Hat Enterprise Linux"
VERSION="8.0 (Ootpa)"
ID="rhel"
ID_LIKE="fedora"
VERSION_ID="8.0"
PLATFORM_ID="platform:el8"
PRETTY_NAME="Red Hat Enterprise Linux 8.0 (Ootpa)"
ANSI_COLOR="0;31"
CPE_NAME="cpe:/o:redhat:enterprise_linux:8.0:latest"
HOME_URL="https://www.redhat.com/"
BUG_REPORT_URL="https://bugzilla.redhat.com/"

REDHAT_BUGZILLA_PRODUCT="Red Hat Enterprise Linux 8"
REDHAT_BUGZILLA_PRODUCT_VERSION=8.0
REDHAT_SUPPORT_PRODUCT="Red Hat Enterprise Linux"
REDHAT_SUPPORT_PRODUCT_VERSION="8.0"
...
```

EXAMPLE #2 (View the Dockerfile in the container) This is another example of running a quick command to inspect the content of a container from the host. All layered images that Red Hat provides include the Dockerfile from which they are built in **/root/buildinfo**. In this case you do not need to mount any volumes from the host.

```
# podman run --rm \
    registry.access.redhat.com/rhel8/rsyslog \
    ls /root/buildinfo
Dockerfile-rhel8-rsyslog-8
```

Now you know what the Dockerfile is called, you can list its contents:

```
# podman run --rm registry.access.redhat.com/rhel8/rsyslog \
  cat /root/buildinfo/Dockerfile-rhel8-rsyslog-8
FROM sha256:eb205f07ce7d0bb63bfe560...
LABEL maintainer="Red Hat, Inc."

RUN INSTALL_PKGS="\
rsyslog \
rsyslog-gnutls \
rsyslog-gssapi \
rsyslog-mysql \
rsyslog-pgsql \
rsyslog-relp \
" && yum -y install $INSTALL_PKGS && rpm -V --nosize
--nofiledigest --nomtime --nomode $INSTALL_PKGS && yum clean all
LABEL com.redhat.component="rsyslog-container"
LABEL name="rhel8/rsyslog"
LABEL version="8.0"
...
```

EXAMPLE #3 (Run a shell inside the container) Using a container to launch a bash shell lets you look inside the container and change the contents. This sets the name of the container to **mybash**. The **-i** creates an interactive session and **-t** opens a terminal session. Without **-i**, the shell would open and then exit. Without **-t**, the shell would stay open, but you wouldn't be able to type anything to the shell.

Once you run the command, you are presented with a shell prompt and you can start running commands from inside the container:

```
# podman run --name=mybash -it registry.redhat.io/ubi8/ubi /bin/bash
[root@ed904b8f2d5c/]# yum install procps-ng
[root@ed904b8f2d5c/]# ps -ef
UID      PID  PPID  C STIME TTY      TIME CMD
root      1    0  0 00:46 pts/0    00:00:00 /bin/bash
root     35    1  0 00:51 pts/0    00:00:00 ps -ef
[root@49830c4f9cc4/]# exit
```

Although the container is no longer running once you exit, the container still exists with the new software package still installed. Use **podman ps -a** to list the container:

```
# podman ps -a
CONTAINER ID IMAGE                COMMAND                  CREATED        STATUS        PORTS NAMES
IS INFRA
1ca061b47bd7 .../ubi8/ubi:latest /bin/bash 8 minutes ago Exited 12 seconds ago musing_brown
false
...
```

You could start that container again using **podman start** with the **-ai** options. For example:

```
# podman start -ai mybash
[root@ed904b8f2d5c/]#
```

EXAMPLE #4 (Bind mounting log files) One way to make log messages from inside a container available to the host system is to bind mount the host's `/dev/log` device inside the container. This example illustrates how to run an application in a RHEL container that is named **log_test** that generates

log messages (just the logger command in this case) and directs those messages to the `/dev/log` device that is mounted in the container from the host. The `--rm` option removes the container after it runs.

```
# podman run --name="log_test" -v /dev/log:/dev/log --rm \
    registry.redhat.io/ubi8/ubi logger "Testing logging to the host"
# journalctl -b | grep Testing
Nov 12 20:00:10 ubi8 root[17210]: Testing logging to the host
```

EXAMPLE #5 (Run a service as a daemon with a static IP address) The following example runs an **rsyslog** service as a daemon process, so it runs continuously in the background. It also tells **podman** to set the container network interface to a particular IP address (for example, 10.88.0.44). After that, you can run **podman inspect** command to check that the IP address was set properly:

```
# podman run -d --ip=10.88.0.44 registry.access.redhat.com/rhel7/rsyslog
efde5f0a8c723f70dd5cb5dc3d5039df3b962fae65575b08662e0d5b5f9fbe85
# podman inspect efde5f0a8c723 | grep 10.88.0.44
    "IPAddress": "10.88.0.44",
```

3.2. INVESTIGATING RUNNING AND STOPPED CONTAINERS

After you have some running containers, you can list both those containers that are still running and those that have exited or stopped with the **podman ps** command. You can also use the **podman inspect** to look at specific pieces of information within those containers.

3.2.1. Listing containers

Let's say you have one or more containers running on your host. To work with containers from the host system, you can open a shell and try some of the following commands.

podman ps: The `ps` option shows all containers that are currently running:

```
# podman run -d registry.redhat.io/rhel8/rsyslog
# podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
74b1da000a11	rhel8/rsyslog	/bin/rsyslog.sh	2 minutes ago	Up	About a minute	musng_brown

If there are containers that are not running, but were not removed (`--rm` option), the containers are still hanging around and can be restarted. The **podman ps -a** command shows all containers, running or stopped.

```
# podman ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES	IS
INFRA							
d65aecc325a4	ubi8/ubi	/bin/bash	3 secs ago	Exited (0)	5 secs ago	peaceful_hopper	false
74b1da000a11	rhel8/rsyslog	rsyslog.sh	2 mins ago	Up	About a minute	musng_brown	false

3.2.2. Inspecting containers

To inspect the metadata of an existing container, use the **podman inspect** command. You can show all metadata or just selected metadata for the container. For example, to show all metadata for a selected container, type:

```
# podman inspect 74b1da000a11
...
"ID": "74b1da000a114015886c557deec8bed9dfb80c888097aa83f30ca4074ff55fb2",
"Created": "2018-11-13T10:30:31.884673073-05:00",
"Path": "/bin/rsyslog.sh",
"Args": [
    "/bin/rsyslog.sh"
],
"State": {
    OciVersion: "1.0.1-dev",
    Status: "running",
    Running: true,
    ...
}
```

You can also use `inspect` to pull out particular pieces of information from a container. The information is stored in a hierarchy. So to see the container's IP address (IPAddress under NetworkSettings), use the **--format** option and the identity of the container. For example:

```
# podman inspect --format='{{.NetworkSettings.IPAddress}}' 74b1da000a11
10.88.0.31
```

Examples of other pieces of information you might want to inspect include `.Path` (to see the command run with the container), `.Args` (arguments to the command), `.Config.ExposedPorts` (TCP or UDP ports exposed from the container), `.State.Pid` (to see the process id of the container) and `.HostConfig.PortBindings` (port mapping from container to host). Here's an example of `.State.Pid` and `.State.StartedAt`:

```
# podman inspect --format='{{.State.Pid}}' 74b1da000a11
19593
# ps -ef | grep 19593
root  19593 19583  0 10:30 ?        00:00:00 /usr/sbin/rsyslogd -n
# podman inspect --format='{{.State.StartedAt}}' 74b1da000a11
2018-11-13 10:30:35.358175255 -0500 EST
```

In the first example, you can see the process ID of the containerized executable on the host system (PID 19593). The **ps -ef** command confirms that it is the `rsyslogd` daemon running. The second example shows the date and time that the container was run.

3.2.3. Investigating within a container

To investigate within a running container, you can use the **podman exec** command. With **podman exec**, you can run a command (such as **/bin/bash**) to enter a running container process to investigate that container.

The reason for using **podman exec**, instead of just launching the container into a bash shell, is that you can investigate the container as it is running its intended application. By attaching to the container as it is performing its intended task, you get a better view of what the container actually does, without necessarily interrupting the container's activity.

Here is an example using **podman exec** to look into a running `rsyslog`, then look around inside that container.

1. **Launch a container:** Launch a container such the `rsyslog` container image described earlier. Type **podman ps** to make sure it is running:

```
# podman ps
CONTAINER ID  IMAGE          COMMAND          CREATED    STATUS    PORTS
NAMES
74b1da000a11  rsyslog:latest "/usr/rsyslog.sh 6 minutes ago Up 6 minutes    rsyslog
```

2. Enter the container with **podman exec**: Use the container ID or name to open a bash shell to access the running container. Then you can investigate the attributes of the container as follows:

```
# podman exec -it 74b1da000a11 /bin/bash
[root@74b1da000a11 /]# cat /etc/redhat-release
Red Hat Enterprise Linux release 8.0
[root@74b1da000a11 /]# yum install procps-ng
[root@74b1da000a11 /]# ps -ef
UID      PID  PPID  C STIME TTY      TIME CMD
root      1    0  0 15:30 ?      00:00:00 /usr/sbin/rsyslogd -n
root      8    0  6 16:01 pts/0    00:00:00 /bin/bash
root     21    8  0 16:01 pts/0    00:00:00 ps -ef
[root@74b1da000a11 /]# df -h
Filesystem      Size  Used Avail Use% Mounted on
overlay          39G   2.5G   37G   7% /
tmpfs            64M    0   64M   0% /dev
tmpfs            1.5G   8.7M   1.5G   1% /etc/hosts
shm              63M    0   63M   0% /dev/shm
tmpfs            1.5G    0   1.5G   0% /sys/fs/cgroup
tmpfs            1.5G    0   1.5G   0% /proc/acpi
tmpfs            1.5G    0   1.5G   0% /proc/scsi
tmpfs            1.5G    0   1.5G   0% /sys/firmware
[root@74b1da000a11 /]# uname -r
4.18.0-80.1.2.el8_0.x86_64
[root@74b1da000a11 /]# rpm -qa | more
redhat-release-8.0-0.44.el8.x86_64
filesystem-3.8-2.el8.x86_64
basesystem-11-5.el8.noarch
ncurses-base-6.1-7.20180224.el8.noarch
...
bash-4.2# free -m
              total        used        free      shared  buff/cache   available
Mem:           1941         560          139          10        1241        1189
Swap:           1023           15         1008
[root@74b1da000a11 /]# exit
```

The commands just run from the bash shell (running inside the container) show you several things.

- The container was built from a RHEL release 8.0 image.
- The process table (`ps -ef`) shows that the `/usr/sbin/rsyslogd` command is process ID 1.
- Processes running in the host's process table cannot be seen from within the container. Although the `rsyslogd` process can be seen on the host process table (it was process ID 19593 on the host).
- There is no separate kernel running in the container (`uname -r` shows the host system's kernel).

- The `rpm -qa` command lets you see the RPM packages that are included inside the container. In other words, there is an RPM database inside of the container.
- Viewing memory (`free -m`) shows the available memory on the host (although what the container can actually use can be limited using `cgroups`).

3.3. STARTING AND STOPPING CONTAINERS

If you ran a container, but didn't remove it (`--rm`), that container is stored on your local system and ready to run again. To start a previously run container that wasn't removed, use the **start** option. To stop a running container, use the **stop** option.

3.3.1. Starting containers

A container that doesn't need to run interactively can sometimes be restarted after being stopped with only the **start** option and the container ID or name. For example:

```
# podman start myrhel_httpd
myrhel_httpd
```

To start a container so you can work with it from the local shell, use the `-a` (attach) and `-i` (interactive) options. Once the bash shell starts, run the commands you want inside the container and type `exit` to kill the shell and stop the container.

```
# podman start -a -i agitated_hopper
[root@d65aecc325a4 /]# exit
```

3.3.2. Stopping containers

To stop a running container that is not attached to a terminal session, use the `stop` option and the container ID or number. For example:

```
# podman stop 74b1da000a11
74b1da000a114015886c557deec8bed9dfb80c888097aa83f30ca4074ff55fb2
```

The **stop** option sends a `SIGTERM` signal to terminate a running container. If the container doesn't stop after a grace period (10 seconds by default), **podman** sends a `SIGKILL` signal. You could also use the **podman kill** command to kill a container (`SIGKILL`) or send a different signal to a container. Here's an example of sending a `SIGHUP` signal to a container (if supported by the application, a `SIGHUP` causes the application to re-read its configuration files):

```
# podman kill --signal="SIGHUP" 74b1da000a11
74b1da000a114015886c557deec8bed9dfb80c888097aa83f30ca4074ff55fb2
```

3.4. REMOVING CONTAINERS

To see a list of containers that are still hanging around your system, run the **podman ps -a** command. To remove containers you no longer need, use the **podman rm** command, with the container ID or name as an option. You should stop any containers that are still running before removing them. Here is an example:

```
# podman rm goofy_wozniak
```

You can remove multiple containers on the same command line:

```
# podman rm clever_yonath furious_shockley drunk_newton
```

If you want to clear out all your containers, you could use a command like the following to remove all containers (not images) from your local system (make sure you mean it before you do this!):

```
# podman rm -a  
56c496350bd534da7620fe2fa660526a6fc7f1c57b0298291cd2210311fe723b  
83ad58c17b20f9e8271171f3023ae094dbfab6ce5708344a68feb121916961ca  
a93b696a1f5629300382a8ce860c4ba42f664db98101e82c2dbcc2074b428faf  
bee71e61b53bd8b036b2e8cb8f570ef8308403502760a27ee23a4b675d92b93d
```

CHAPTER 4. USING RED HAT UNIVERSAL BASE IMAGES (STANDARD, MINIMAL, AND RUNTIMES)

Red Hat Enterprise Linux (RHEL) base images are meant to form the foundation for the container images you build. For RHEL 8, all Red Hat base images are available as new Universal Base Images (UBI). These include versions of RHEL standard, minimal, init, and Red Hat Software Collections that are all now freely available and redistributable. Characteristics of RHEL base images include:

- **Supported:** Supported by Red Hat for use with your containerized applications. Contains the same secured, tested, and certified software packages you have in Red Hat Enterprise Linux.
- **Cataloged:** Listed in the [Red Hat Container Catalog](#), where you can find descriptions, technical details, and a health index for each image.
- **Updated:** Offered with a well-defined update schedule, so you know you are getting the latest software (see [Red Hat Container Image Updates](#)).
- **Tracked:** Tracked by errata, to help you understand the changes that go into each update.
- **Reusable:** Only need to be downloaded and cached in your production environment once, where each base image can be reused by all containers that include it as their foundation.

Red Hat Universal Base Images (UBI) for RHEL 8 provide the same quality RHEL software for building container images as their non-UBI predecessors (**rhel6**, **rhel7**, **rhel-init**, and **rhel-minimal** base images), but offer more freedom in how they are used and distributed.



NOTE

For a list of available Red Hat UBI images, and associated information about UBI repositories and source code, see [Universal Base Images \(UBI\): Images, repositories, and packages](#).

4.1. WHAT ARE RED HAT BASE IMAGES?

Red Hat provides multiple base images that you can use as a starting point for your own images. These images are available through the Red Hat Registry (registry.access.redhat.com and registry.redhat.io) and described in the [Red Hat Container Catalog](#).

For RHEL 8, there are standard, minimal and init base image available. Red Hat also provides a set of language runtime images, based on [Application Streams](#), that you can build on when you are creating containers for applications that require specific runtimes. Runtime images include python, php, ruby, nodejs, and others. All of the RHEL 8 images are UBI images, which means that you can freely obtain and redistribute them.

There is a set of RHEL 7 images as well that you can run on RHEL 8 systems. For RHEL 7, there are both UBI (redistributable) and non-UBI (require subscription access and are non-redistributable) base images. Those images include three regular base images (**rhel7**, **rhel-init**, and **rhel-minimal**) and three UBI images (**ubi7**, **ubi7-init**, and **ubi7-minimal**).

Although Red Hat does not offer tools for running containers on RHEL 6 systems, it does offer RHEL 6 container images you can use. There are standard (**rhel6**) and Init (**rhel6-init**) base image available for RHEL 6, but no minimal RHEL 6 image. Likewise, there are no RHEL 6 UBI images.

4.1.1. Using standard Red Hat base images

Standard RHEL 8 base images (**ubi8**) have a robust set of software features that include the following:

- **init system:** All the features of the systemd initialization system you need to manage systemd services are available in the standard base images. These init systems let you install RPM packages that are pre-configured to start up services automatically, such as a Web server (httpd) or FTP server (vsftpd).
- **yum:** Software needed to install software packages is included via the standard set of **yum** commands (**yum**, **yum-config-manager**, **yumdownloader**, and so on). For the UBI base images, you have access to free yum repositories for adding and updating software.
- **utilities:** The standard base image includes some useful utilities for working inside the container. Utilities that are in this base image that are not in the minimal images include **tar**, **dmidecode**, **gzip**, **getfacl** (and other acl commands), **dmsetup** (and other device mapper commands), and others.

4.1.2. Using minimal Red Hat base images

The **ubi8-minimal** images are stripped-down RHEL images to use when a bare-bones base image is desired. If you are looking for the smallest possible base image to use as part of the larger Red Hat ecosystem, you can start with these minimal images.

RHEL minimal images provide a base for your own container images that is less than half the size of the standard image, while still being able to draw on RHEL software repositories and maintain any compliance requirements your software has.

Here are some features of the minimal base images:

- **Small size:** Minimal images are about 92M on disk and 32M compressed. This makes it less than half the size of the standard images.
- **Software installation (microdnf):** Instead of including the full-blown yum facility for working with software repositories and RPM software packages, the minimal images includes the microdnf utility. Microdnf is a scaled-down version of dnf. It includes only what is needed to enable and disable repositories, as well as install, remove, and update packages. It also has a clean option, to clean out cache after packages have been installed.
- **Based on RHEL packaging:** Because minimal images incorporate regular RHEL software RPM packages, with a few features removed such as extra language files or documentation, you can continue to rely on RHEL repositories for building your images. This allows you to still maintain compliance requirements you have that are based on RHEL software. Features of minimal images make them perfect for trying out applications you want to run with RHEL, while carrying the smallest possible amount of overhead. What you don't get with minimal images is an initialization and service management system (systemd or System V init), a Python run-time environment, and a bunch of common shell utilities.
- **Modules for microdnf are not supported:** Modules used with the **dnf** command let you install multiple versions of the same software, when available. The **microdnf** utility included with minimal images does not support modules. So if modules are required, you should use a non-minimal base images, which include yum.

If your goal, however, is just to try to run some simple binaries or pre-packaged software that doesn't have a lot of requirements from the operating system, the minimal images might suit your needs. If your application does have dependencies on other software from RHEL, you can simply use microdnf to install the needed packages at build time.

Red Hat intends for you to always use the latest version of the minimal images, which is implied by simply requesting **ubi8/ubi-minimal** or **ubi8-minimal**. Red Hat does not expect to support older versions of minimal images going forward.

4.1.3. Using Init Red Hat base images

The UBI **ubi8-init** images contains the systemd initialization system, making them useful for building images in which you want to run systemd services, such as a web server or file server. The Init image contents are less than what you get with the standard images, but more than what is in the minimal images.



NOTE

Because the **ubi8-init** image builds on top of the **ubi8** image, their contents are mostly the same. There are a few critical differences, however. In **ubi8-init**, the `Cmd` is set to `/sbin/init`, instead of `bash`, to start the systemd Init service by default. It includes `ps` and process related commands (procpss-ng package), which **ubi8** does not. Also, **ubi8-init** sets `SIGRTMIN+3` as the `StopSignal`, as systemd in **ubi8-init** ignores normal signals to exit (`SIGTERM` and `SIGKILL`), but will terminate if it receives `SIGRTMIN+3`.

Historically, Red Hat Enterprise Linux base container images were designed for Red Hat customers to run enterprise applications, but were not free to redistribute. This can create challenges for some organizations that need to redistribute their applications. That's where the Red Hat Universal Base Images come in.

4.2. HOW ARE UBI IMAGES DIFFERENT?

UBI images were created so you can build your container images on a foundation of official Red Hat software that can be freely shared and deployed. From a technical perspective, they are nearly identical to legacy Red Hat Enterprise Linux images, which means they have great security, performance, and life cycles, but they are released under a different End User License Agreement. Here are some attributes of Red Hat UBI images:

- **Built from a subset of RHEL content** Red Hat Universal Base images are built from a subset of normal Red Hat Enterprise Linux content. All of the content used to build selected UBI images is released in a publicly available set of yum repositories. This lets you install extra packages, as well as update any package in UBI base images.
- **Redistributable:** The intent of UBI images is to allow Red Hat customers, partners, ISVs, and others to standardize on one container base image, allowing users to focus on application needs instead of distribution rules. These images can be shared and run in any environment capable of running those images. As long as you follow some basic guidelines, you will be able to freely redistribute your UBI-based images.
- **Base and runtime images** Besides the three types of base images, UBI versions of various runtime images are available as well. These runtime images provide a foundation for applications that can benefit from standard, supported runtimes such as python, php, nodejs, and ruby.
- **Enabled yum repositories:** The following yum repositories are enabled within each RHEL 8 UBI image:
 - The **ubi-8-baseos** repo holds the redistributable subset of RHEL packages you can include in your container.

- The **ubi-8-appstream** repo holds Red Hat Software Collections packages that you can add to a UBI image to help you standardize the environments you use with applications that require particular runtimes.
- **Licensing:** You are free to use and redistribute UBI images, provided you adhere to the [Red Hat Universal Base Image End User Licensing Agreement](#).
- **Adding UBI RPMs:** You can add RPM packages to UBI images from preconfigured UBI repositories. If you happen to be in a disconnected environment, you must whitelist the UBI Content Delivery Network (<https://cdn-ubi.redhat.com>) to use that feature. See the [Connect to https://cdn-ubi.redhat.com](https://cdn-ubi.redhat.com) solution for details.

Although the legacy RHEL 7 base images will continue to be supported, UBI images are recommended going forward. For that reason, examples in the rest of this chapter are done with RHEL 8 UBI images.

4.3. GET UBI IMAGES

For more information about available Red Hat Universal Base Images, see the article [Universal Base Images \(UBI\): Images, repositories, packages, and source code](#).

4.4. PULL UBI IMAGES

To pull UBI images to your system so you can use them with tools such as podman, buildah or skopeo, type the following:

```
# podman pull registry.access.redhat.com/ubi8/ubi:latest
# podman pull registry.access.redhat.com/ubi8/ubi-minimal:latest
```

To check that the images are available on your system, type:

```
# podman images
REPOSITORY                                TAG  IMAGE ID    CREATED    SIZE
registry.access.redhat.com/ubi8/ubi-minimal latest c94a444803e3 8 hours ago 80.9 MB
registry.access.redhat.com/ubi8/ubi      latest 40b488f87628 17 hours ago 214 MB
```

When pulled in this way, images are available and usable by **podman**, **buildah**, **skopeo** and the CRI-O container image, but they are not available to the Docker service or **docker** command. To use these images with Docker, you can run **docker pull** instead.

4.5. REDISTRIBUTING UBI IMAGES

After you pull a UBI image, you are free to push it to your own registry and share it with others. You can upgrade or add to that image from UBI yum repositories as you like. Here is an example of how to push a UBI image to your own or another third-party repository:

```
# podman pull registry.redhat.io/ubi8/ubi
# podman tag registry.access.redhat.com/ubi8/ubi registry.example.com:5000/ubi8/ubi
# podman push registry.example.com:5000/ubi8/ubi
```

While there are few restrictions on how you use these images, there are some restrictions about how you can refer to them. For example, you can't call those images Red Hat certified or Red Hat supported unless you certify it through the [Red Hat Partner Connect Program](#), either with Red Hat Container Certification or Red Hat OpenShift Operator Certification.

4.6. RUN UBI IMAGES

To start a container from a UBI image and run the bash shell in that image (so you can look around inside), do the following (type exit when you are done):

```
# podman run --rm -it registry.access.redhat.com/ubi8/ubi-minimal:latest /bin/bash
[root@da9213157c51 /]#
# podman run --rm -it registry.access.redhat.com/ubi8/ubi:latest /bin/bash
bash-4.2#
```

While in the container:

- Run **rpm -qa** to see a list of package inside each container.
- Type **yum list available** to see packages available to add to the image from the UBI yum repos. (The yum command is not available in the **ubi-minimal** containers.)
- Get source code, as described in the "Getting UBI Container Image Source Code," later in this chapter.

On systems that include the Docker service, you can use **docker run** instead.

4.7. ADD SOFTWARE TO A RUNNING UBI CONTAINER

UBI images are built from 100% Red Hat content. These UBI images also provide a subset of Red Hat Enterprise Linux packages that are freely available to install for use with UBI. To add or update software, UBI images are pre-configured to point to the freely available yum repositories that hold official Red Hat RPMs.

To add packages from UBI repos to running UBI containers:

- On **ubi** images, the yum command is installed to let you draw packages
- On **ubi-minimal** images, the **microdnf** command (with a smaller feature set) is included instead of **yum**.

Keep in mind that installing and working with software packages directly in running containers is just for adding packages temporarily or learning about the repos. Refer to the "Build a UBI-based image" for more permanent ways of building UBI-based images.

When you add software to a UBI container, procedures differ for updating UBI images on a subscribed RHEL host or on an unsubscribed (or non-RHEL) system. Those two ways of working with UBI images are illustrated below.

4.7.1. Adding software to a UBI container (subscribed host)

If you are running a UBI container on a registered and subscribed RHEL host, the main RHEL Server repository is enabled inside the standard UBI container, along with all the UBI repos. So the full set of Red Hat packages is available. From the UBI minimal container, all UBI repos are enabled by default, but no repos are enabled from the host by default.

4.7.2. Adding software inside the standard UBI container

To ensure the containers you build can be redistributed, disable non-UBI yum repositories in the standard UBI image when you add software. If you disable all yum repositories except for UBI repositories, only packages from the freely available repos are used when you add software.

With a shell open inside a standard UBI base image container (**ubi8/ubi**) from a subscribed RHEL host, run the following command to add a package to that container (for example, the `bzip2` package):

```
# yum install --disablerepo=* --enablerepo=ubi-8-appstream --enablerepo=ubi-8-baseos bzip2
```

To add software inside a standard UBI container that is in the RHEL server repo, but not in UBI repos, don't disable any repositories and just install the package:

```
# yum install zsh
```

To install a package that is in a different host repo from inside the standard UBI container, you have to explicitly enable the repo you need. For example:

```
# yum install --enablerepo=rhel-7-server-optional-rpms zsh-html
```



WARNING

Installing Red Hat packages that are not inside the Red Hat UBI repos might limit how widely you can distribute the container outside of subscribed hosts.

4.7.3. Adding software inside the minimal UBI container

UBI yum repositories are enabled inside the UBI minimal image by default.

To install the same package demonstrated earlier (`bzip2`) from one of those UBI yum repositories on a subscribed RHEL host from the UBI minimal container, type:

```
# microdnf install bzip2
```

To install packages inside a minimal UBI container from repos available on a subscribed host that are not part of a UBI yum repo, you would have to explicitly enable those repos. For example:

```
# microdnf install --enablerepo=rhel-7-server-rpms zsh
# microdnf install --enablerepo=rhel-7-server-rpms \
    --enablerepo=rhel-7-server-optional-rpms zsh-html
```

**WARNING**

Using non-UBI RHEL repositories to install packages in your UBI images could restrict your ability to share those images to run outside of subscribed RHEL systems.

4.7.4. Adding software to a UBI container (unsubscribed host)

To add software packages to a running container that is either on an unsubscribed RHEL host or some other Linux system, you don't have to disable any yum repositories. For example:

```
# yum install bzip2
```

To install that package on an unsubscribed RHEL host from the UBI minimal container, type:


```
# microdnf install bzip2
```

As noted earlier, both of these means of adding software to a running UBI container are not intended for creating permanent UBI-based container images. For that, you should build new layers on to UBI images, as described in the following section.

4.7.5. Build a UBI-based image

You can build UBI-based container images in the same way you build other images, with one exception. You should disable all non-UBI yum repositories when you actually build the images, if you want to be sure that your image only contains Red Hat software that you can redistribute.

Here's an example of creating a UBI-based Web server container from a Dockerfile with the **buildah** utility:

**NOTE**

For ubi8/ubi-minimal images, use microdnf instead of yum below:

```
RUN microdnf update -y && rm -rf /var/cache/yum
RUN microdnf install httpd -y && microdnf clean all
```

1. **Create a Dockerfile** Add a **Dockerfile** with the following contents to a new directory:

```
FROM registry.access.redhat.com/ubi8/ubi
USER root
LABEL maintainer="John Doe"
# Update image
RUN yum update --disablerepo=* --enablerepo=ubi-8-appstream --enablerepo=ubi-8-baseos
-y && rm -rf /var/cache/yum
RUN yum install --disablerepo=* --enablerepo=ubi-8-appstream --enablerepo=ubi-8-baseos
httpd -y && rm -rf /var/cache/yum
# Add default Web page and expose port
RUN echo "The Web Server is Running" > /var/www/html/index.html
```

```
EXPOSE 80
# Start the service
CMD ["-D", "FOREGROUND"]
ENTRYPOINT ["/usr/sbin/httpd"]
```

2. **Build the new image** While in that directory, use **buildah** to create a new UBI layered image:

```
# buildah bud -t johndoe/webserver .
STEP 1: FROM registry.access.redhat.com/ubi8/ubi:latest
STEP 2: USER root
STEP 3: LABEL maintainer="John Doe"
STEP 4: RUN yum update --disablerepo=* --enablerepo=ubi-8-appstream --enablerepo=ubi-8-baseos -y
...
No packages marked for update
STEP 5: RUN yum install --disablerepo=* --enablerepo=ubi-8-appstream --enablerepo=ubi-8-baseos httpd -y
Loaded plugins: ovl, product-id, search-disabled-repos
Resolving Dependencies
--> Running transaction check
=====
Package                Arch      Version              Repository            Size
=====
Installing:
httpd                  x86_64    2.4.37-10            latest-rhubi-8.0-appstream 1.4 M
Installing dependencies:
apr                    x86_64    1.6.3-9.el8          latest-rhubi-8.0-appstream 125 k
apr-util               x86_64    1.6.1-6.el8          latest-rhubi-8.0-appstream 105 k
httpd-filesystem       noarch    2.4.37-10            latest-rhubi-8.0-appstream 34 k
httpd-tools            x86_64    2.4.37-10.
...

Transaction Summary
...
Complete!
STEP 6: RUN echo "The Web Server is Running" > /var/www/html/index.html
STEP 7: EXPOSE 80
STEP 8: CMD ["-D", "FOREGROUND"]
STEP 9: ENTRYPOINT ["/usr/sbin/httpd"]
STEP 10: COMMIT
...
Writing manifest to image destination
Storing signatures
--> 36a604cc0dd3657b46f8762d7ef69873f65e16343b54c63096e636c80f0d68c7
```

3. **Test:** Test the UBI layered webserver image:

```
# podman run -d -p 80:80 johndoe/webserver
bbe98c71d18720d966e4567949888dc4fb86eec7d304e785d5177168a5965f64
# curl http://localhost/index.html
The Web Server is Running
```

4.7.6. Using AppStream runtime images

Red Hat Enterprise Linux 8 AppStream offers another set of container images that you can use as the basis for your container builds. These images are built on RHEL standard base images, with most already updated as UBI images. Each of these images include additional software you might want to use for specific runtime environments.

So, if you expect to build multiple images that require, for example, php runtime software, you can use provide a more consistent platform for those images by starting with a PHP software collections image.

Here are a few examples of AppStream container images built on UBI base images, that are available from the Red Hat Registry (registry.access.redhat.com or registry.redhat.io):

- **ubi8/php-72**: PHP 7.2 platform for building and running applications
- **ubi8/nodejs-10**: Node.js 10 platform for building and running applications. Used by Node.js 10 Source-To-Image builds
- **ubi8/ruby25**: Ruby 2.5 platform for building and running applications
- **ubi8/python-27**: Python 2.7 platform for building and running applications
- **ubi8/python-36**: Python 3.6 platform for building and running applications
- **ubi8/s2i-core**: Base image with essential libraries and tools used as a base for builder images like perl, python, ruby, and so on
- **ubi8/s2i-base**: Base image for Source-to-Image builds

Because these UBI images contain the same basic software as their legacy image counterparts, you can learn about those images from the [Using Red Hat Software Collections Container Images](#) guide. Be sure to use the UBI image names to pull those images.

RHEL 8 AppStream container images are updated every time RHEL 8 base images are updated. For RHEL 7, these same images (referred to as Red Hat Software Collections images) are updated on a schedule that is separate from RHEL base image updates (as are related images for Dotnet and DevTools). Search the [Red Hat Container Catalog](#) for details on any of these images. For more information on update schedules, see [Red Hat Container Image Updates](#).

4.7.7. Getting UBI Container Image Source Code

Source code is available for all Red Hat UBI-based images in the form of downloadable containers. Before continuing, be aware about Red Hat source containers:

- Source container images cannot be run, despite being packaged as containers. To install Red Hat source container images on your system, use the **skopeo command**, instead of using **podman pull** command.
 - Use **skopeo copy** command to copy a source container image to a directory on your local system.
 - Use **skopeo inspect** command to inspect the source container image.
- For more details on **skopeo** command, see [Section 1.5. Using skopeo to work with container registries](#).
- Source container images are named based on the binary containers they represent. For

example, for a particular standard RHEL UBI 8 container **registry.access.redhat.com/ubi8:8.1-397** append **-source** to get the source container image (**registry.access.redhat.com/ubi8:8.1-397-source**).

- Once a source container image is copied to a local directory, you can use a combination of **tar**, **gzip**, and **rpm** commands to work with that content.
- It could take several hours after a container image is released for its associated source container to become available.

Procedure

1. Use **skopeo copy** command to copy the source container image to a local directory:

```
$ skopeo copy \
docker://registry.access.redhat.com/ubi8:8.1-397-source \
dir:$HOME/TEST
...
Copying blob 477bc8106765 done
Copying blob c438818481d3 done
Copying blob 26fe858c966c done
Copying blob ba4b5f020b99 done
Copying blob f7d970ccd456 done
Copying blob ade06f94b556 done
Copying blob cc56c782b513 done
Copying blob dcf9396fdada done
Copying blob feb6d2ae2524 done
Copying config dd4cd669a4 done
Writing manifest to image destination
Storing signatures
```

2. Use **skopeo inspect** command to inspect the source container image:

```
$ skopeo inspect dir:$HOME/TEST
{
  "Digest":
"sha256:7ab721ef3305271bbb629a6db065c59bbeb87bc53e7cbf88e2953a1217ba7322",
  "RepoTags": [],
  "Created": "2020-02-11T12:14:18.612461174Z",
  "Dockerversion": "",
  "Labels": null,
  "Architecture": "amd64",
  "Os": "linux",
  "Layers": [
    "sha256:1ae73d938ab9f11718d0f6a4148eb07d38ac1c0a70b1d03e751de8bf3c2c87fa",
    "sha256:9fe966885cb8712c47efe5ecc2eaa0797a0d5ffb8b119c4bd4b400cc9e255421",
    "sha256:61b2527a4b836a4efbb82dfd449c0556c0f769570a6c02e112f88f8bbcd90166",
    ...
    "sha256:cc56c782b513e2bdd2cc2af77b69e13df4ab624ddb856c4d086206b46b9b9e5f",
    "sha256:dcf9396fdada4e6c1ce667b306b7f08a83c9e6b39d0955c481b8ea5b2a465b32",

"sha256:feb6d2ae252402ea6a6fca8a158a7d32c7e4572db0e6e5a5eab15d4e0777951e"
  ],
  "Env": null
}
```

3. To untar all the content, type:

```
$ cd $HOME/TEST
$ for f in $(ls); do tar xvf $f; done
```

4. To check the results, type:

```
$ find blobs/ rpm_dir/
blobs/
blobs/sha256
blobs/sha256/10914f1fff060ce31388f5ab963871870535aaaa551629f5ad182384d60fdf82
rpm_dir/
rpm_dir/gzip-1.9-4.el8.src.rpm
```

5. Begin examining and using the content.

4.7.8. Tips and tricks for using UBI images

Here are a few issues to consider when working with UBI images:

- Hundreds of RPM packages used in existing Red Hat Software Collections runtime images are stored in the yum repositories packaged with the new UBI images. Feel free to install those RPMs on your UBI images to emulate the runtime (python, php, nodejs, etc.) that interests you.
- Because some language files and documentation have been stripped out of the minimal UBI image (**ubi8/ubi-minimal**), running **rpm -Va** inside that container will show the contents of many packages as being missing or modified. If having a complete list of files inside that container is important to you, consider using a tool such as **Tripwire** to record the files in the container and check it later.
- After a layered image has been created, use **podman history** to check which UBI image it was built on. For example, after completing the webserver example shown earlier, type **podman history johndoe/webserver** to see that the image it was built on includes the image ID of the UBI image you added on the FROM line of the Dockerfile.

4.7.9. How to request new features in UBI?

Red Hat partners and customers can request new features, including package requests, by filing a support ticket through standard methods. Non-Red Hat customers do not receive support, but can file requests through the standard Red Hat Bugzilla for the appropriate RHEL product. See also: [Red Hat Bugzilla Queue](#)

4.7.10. How to file a support case for UBI?

Red Hat partners and customers can file support tickets through standard methods when running UBI on a supported Red Hat platform (OpenShift/RHEL). Red Hat support staff will guide partners and customers

See also: [Open a Support Case](#)

CHAPTER 5. RUNNING SPECIAL CONTAINER IMAGES

Once you become familiar with common ways of working with containers and container images, use this section to learn about some special types of container images you might find useful. These include:

- **Toolbox:** Instead of burdening a host system by installing tools needed to debug problems or monitor features, you can run the **toolbox** command. Toolbox starts a **support-tools** container image that holds tools you can use to run reports or diagnose problems on the host.
- **Runlabels:** Some container images have labels built in that allow you to run those containers with preset options and arguments. When you do a **podman run** with a runlabel, the result can be a prescriptive set of features when you install, run, remove, or upgrade a container image.

5.1. TROUBLESHOOTING CONTAINER HOSTS WITH TOOLBOX

Instead of installing troubleshooting tools directly to your Red Hat Enterprise Linux 8 system, the **toolbox** utility offers a way to temporarily add those tools, then easily discard them when you are done. The **toolbox** utility works by:

- Pulling the **registry.redhat.io/rhel8/support-tools** image to your local system.
- Starting up a container from the image, then running a shell inside the container from which you can access the host system.

The **support-tools** container allows you to:

- Run commands that may not be installed on the host system, such as **sosreport**, **strace**, or **tcpdump**, in a way that lets them act on the host system.
- Install more software inside the container to use on the host system.
- Discard the container when you are done.

The following illustrates a typical **toolbox** session.

Procedure

1. Install the **toolbox** and **podman** packages, if you haven't already. A good way to do that is to install the full set of container tools:

```
# yum module install container-tools -y
```

2. Run the **toolbox** command to pull and run the **support-tools** image (inserting your Red Hat Customer Portal credentials when prompted):

```
# toolbox
Trying to pull registry.redhat.io/rhel8/support-tools...
...
Would you like to authenticate to registry: 'registry.redhat.io' and try again? [y/N] y
Username: johndoe
Password: *****
Login Succeeded!
Trying to pull registry.redhat.io/rhel8/support-tools...Getting image source signatures
...
Storing signatures
```

```
30e261462851238d38f4ef2afdaf55f1f8187775c5ca373b43e0f55722faaf97
Spawning a container 'toolbox-root' with image 'registry.redhat.io/rhel8/support-tools'
Detected RUN label in the container image. Using that as the default...
command: podman run -it --name toolbox-root --privileged --ipc=host --net=host --pid=host -e
HOST=/host -e NAME=toolbox-root -e IMAGE=registry.redhat.io/rhel8/support-tools:latest -v
/run:/run -v /var/log:/var/log -v /etc/machine-id:/etc/machine-id -v /etc/localtime:/etc/localtime -
v /:/host registry.redhat.io/rhel8/support-tools:latest
```

```
bash-4.4#
```

A bash shell opens, ready for you to run commands inside the container.

- From inside the container, the root file system on the host is available from the **/host** directory. The other directories shown are all inside the container.

```
# ls /
bin  dev  home  lib  lost+found  mnt  proc  run  srv  tmp  var
boot  etc  host  lib64  media  opt  root  sbin  sys  usr
```

- From inside the container, you can try a command. For example, you can run **sosreport** to generate information about your system to send to Red Hat support:

```
bash-4.4# sosreport

sosreport (version 3.6)
This command will collect diagnostic and configuration information from
this Red Hat Enterprise Linux system and installed applications.

An archive containing the collected information will be generated in
/host/var/tmp/sos.u82evib and may be provided to a Red Hat support
representative.

...
Press ENTER to continue, or CTRL-C to quit.  <Press ENTER>

...
Your sosreport has been generated and saved in:
  /host/var/tmp/sosreport-rhel81beta-12345678-2019-10-29-pmgjncg.tar.xz
The checksum is: c4e1fd3ee45f78a17afb4e45a05842ed
Please send this file to your support representative.
```

Notice that **sosreport** is aware that you are in a container. As a result it knows to run on the host and save the report to the host (**/host/var/tmp/sosreport-...**).

- Install a software package inside the container, to add tools that are not already in the container. For example, to get a core dump of a running process on the host, install the **procps** and **gcore** packages, use **ps** to get the process ID of a running daemon, then use **gcore** to get a core dump:

```
bash-4.4# yum install procps gdb -y
bash-4.4# ps -ef | grep chronyd
994      809    1 0 Oct28 ?        00:00:00 /usr/sbin/chronyd
bash-4.4# gcore -o /host/tmp/chronyd.core 809
Missing separate debuginfo for target:/usr/sbin/chronyd
Try: dnf --enablerepo='*debug*' install /usr/lib/debug/.build-
id/96/0789a8a3bf28932b093e94b816be379f16a56a.debug
...
```

```

Saved corefile /host/tmp/chronyd.core.809
[Inferior 1 (process 809) detached]
# exit

```

Once you type **exit**, you leave the container and return to the host. You can see that the file saved to **/host/tmp/chronyd.core.809** is available from **/tmp/chronyd.core.809** on the host.

At this point, the container is no longer running, but still exists on the system. You can choose to:

- **Start up the container again** Type **toolbox** again to restart the container (named **toolbox-root**). Any software additions or changes made previously to the container are maintained.
- **Start with a fresh container.** To get rid of the old container, type **podman rm toolbox-root**. Then run **toolbox** again to start with a fresh **support-tools** container.
- **Start with different values** You can change the registry, image, or container name used by **toolbox** by adding the following values to your host **/root/.toolboxrc** file:
 - **REGISTRY:** Change the registry from which the toolbox image is pulled. For example: **REGISTRY=registry.example.com**
 - **IMAGE:** Change the image that is used. For example, **IMAGE=mysupport-tools**
 - **TOOLBOX_NAME:** Change the name assigned to the running container. For example, **TOOLBOX_NAME=mytoolbox**

The next time you run **toolbox**, the new values from the **.toolboxrc** file are used.

5.1.1. Opening privileges to the host

When you run other commands from within the support-tools container (or any privileged container), they can behave differently then when run in a non-privileged container. Although **sosreport** can tell when it is running in a container, other commands would have to specifically be told to act on the host system (the **/host** directory). Here are examples of features that may or not be open to the host from a container:

- **Privileges:** A privileged container (**--privileged**) runs applications as root user on the host by default. The container has this ability because it runs with an **unconfined_t** SELinux security context. So you would be able to, for example, delete files and directories mounted from the host that were owned by the root user.
- **Process tables:** Unlike a regular container, that only sees the processes running inside the container, running a **ps -e** command within a privileged container (with **--pid=host** set) lets you see every process running on the host. So, you can pass a process ID from the host to commands that run in the privileged container (for example, **kill <PID>**). With some commands, however, permissions issues could occur when they try to access processes from the container.
- **Network interfaces:** By default, a container has only one external network interface and one loopback network interface. With network interfaces open to the host (**--net=host**), you can access those network interfaces directly from within the container.
- **Inter-process communications:** The IPC facility on the host is accessible from within the privileged container. So, you can run commands such as **ipcs** to see information about active message queues, shared memory segments, and semaphore sets on the host.

5.2. RUNNING CONTAINERS WITH RUNLABELS

Some Red Hat images include labels that provide pre-set command lines for working with those images. Using the **podman container runlabel <label>** command, you can tell **podman** to execute the command defined in that <label> for the image. Existing runlabels include:

- **install**: Sets up the host system before executing the image. Typically, this results in creating files and directories on the host that the container can access when it is run later.
- **run**: Identifies podman command line options to use when running the container. Typically, the options will open privileges on the host and mount the host content the container needs to remain permanently on the host.
- **uninstall**: Cleans up the host system after you are done running the container.

Red Hat images that have one or more runlabels include the rsyslog and support-tools images. The following procedure illustrates how to use those images.

5.2.1. Running rsyslog with runlabels

The rhel8/rsyslog container image is made to run a containerized version of the rsyslogd daemon. Inside the rsyslog image are **install**, **run** and **uninstall** runlabels. The following procedure steps you through installing, running, and uninstalling the rsyslog image:

Procedure

1. Pull the rsyslog image:

```
# podman pull registry.redhat.io/rhel8/rsyslog
```

2. Display (but do not yet run) the **install** runlabel for rsyslog:

```
# podman container runlabel install --display rhel8/rsyslog
command: podman run --rm --privileged -v /:/host -e HOST=/host -e
IMAGE=registry.redhat.io/rhel8/rsyslog:latest -e NAME=rsyslog
registry.redhat.io/rhel8/rsyslog:latest /bin/install.sh
```

This shows that the command will open privileges to the host, mount the host root filesystem on **/host** in the container, and run an **install.sh** script.

3. Run the **install** runlabel for rsyslog:

```
# podman container runlabel install rhel8/rsyslog
command: podman run --rm --privileged -v /:/host -e HOST=/host -e
IMAGE=registry.redhat.io/rhel8/rsyslog:latest -e NAME=rsyslog
registry.redhat.io/rhel8/rsyslog:latest /bin/install.sh
Creating directory at /host/etc/pki/rsyslog
Creating directory at /host/etc/rsyslog.d
Installing file at /host/etc/rsyslog.conf
Installing file at /host/etc/sysconfig/rsyslog
Installing file at /host/etc/logrotate.d/syslog
```

This creates files on the host system that the rsyslog image will use later.

4. Display the **run** runlabel for rsyslog:

```
# podman container runlabel run --display rhel8/rsyslog
```

```
command: podman run -d --privileged --name rsyslog --net=host --pid=host -v
/etc/pki/rsyslog:/etc/pki/rsyslog -v /etc/rsyslog.conf:/etc/rsyslog.conf -v
/etc/sysconfig/rsyslog:/etc/sysconfig/rsyslog -v /etc/rsyslog.d:/etc/rsyslog.d -v /var/log:/var/log
-v /var/lib/rsyslog:/var/lib/rsyslog -v /run:/run -v /etc/machine-id:/etc/machine-id -v
/etc/localtime:/etc/localtime -e IMAGE=registry.redhat.io/rhel8/rsyslog:latest -e
NAME=rsyslog --restart=always registry.redhat.io/rhel8/rsyslog:latest /bin/rsyslog.sh
```

This shows that the command opens privileges to the host and mount a bunch of files and directories from the host inside the container, when it launches the rsyslog container to run the rsyslogd daemon.

5. Execute the **run** runlabel for rsyslog:

```
# podman container runlabel run rhel8/rsyslog
command: podman run -d --privileged --name rsyslog --net=host --pid=host -v
/etc/pki/rsyslog:/etc/pki/rsyslog -v /etc/rsyslog.conf:/etc/rsyslog.conf -v
/etc/sysconfig/rsyslog:/etc/sysconfig/rsyslog -v /etc/rsyslog.d:/etc/rsyslog.d -v /var/log:/var/log
-v /var/lib/rsyslog:/var/lib/rsyslog -v /run:/run -v /etc/machine-id:/etc/machine-id -v
/etc/localtime:/etc/localtime -e IMAGE=registry.redhat.io/rhel8/rsyslog:latest -e
NAME=rsyslog --restart=always registry.redhat.io/rhel8/rsyslog:latest /bin/rsyslog.sh
28a0d719ff179adcea81eb63cc90fcd09f1755d5edb121399068a4ea59bd0f53
```

The rsyslog container opens privileges, mounts what it needs from the host, and runs the rsyslogd daemon in the background (-d). The rsyslogd daemon begins gathering log messages and directing messages to files in the **/var/log** directory.

6. Display the **uninstall** runlabel for rsyslog:

```
# podman container runlabel uninstall --display rhel8/rsyslog
command: podman run --rm --privileged -v /:/host -e HOST=/host -e
IMAGE=registry.redhat.io/rhel8/rsyslog:latest -e NAME=rsyslog
registry.redhat.io/rhel8/rsyslog:latest /bin/uninstall.sh
```

7. Run the **uninstall** runlabel for rsyslog:

```
# podman container runlabel uninstall rhel8/rsyslog
command: podman run --rm --privileged -v /:/host -e HOST=/host -e
IMAGE=registry.redhat.io/rhel8/rsyslog:latest -e NAME=rsyslog
registry.redhat.io/rhel8/rsyslog:latest /bin/uninstall.sh
```

In this case, the **uninstall.sh** script just removes the **/etc/logrotate.d/syslog** file. Note that it doesn't clean up the configuration files.

5.2.2. Running support-tools with runlabels

The rhel8/support-tools container image is made to run tools such as sosreport and sos-collector to help you analyze your host system. To simplify running the support-tools image, it includes a **run** runlabel. The following procedure describes how to run the support-tools image:

Procedure

1. Pull the support-tools image:

```
# podman pull registry.redhat.io/rhel8/support-tools
```

2. Display (but do not yet run) the **run** runlabel for support-tools:

```
# podman container runlabel run --display rhel8/support-tools
command: podman run -it --name support-tools --privileged --ipc=host --net=host --pid=host -
e HOST=/host -e NAME=support-tools -e IMAGE=registry.redhat.io/rhel8/support-tools:latest
-v /run:/run -v /var/log:/var/log -v /etc/machine-id:/etc/machine-id -v
/etc/localtime:/etc/localtime -v /:/host registry.redhat.io/rhel8/support-tools:latest
```

This shows that the command mounts directories and opens privileges and namespaces (ipc, net, and pid) to the host system. It assigns the host's root file system to the **/host** directory in the container.

3. Execute the **run** runlabel for support-tools:

```
# podman container runlabel run rhel8/support-tools
command: podman run -it --name support-tools --privileged --ipc=host --net=host --pid=host -
e HOST=/host -e NAME=support-tools -e IMAGE=registry.redhat.io/rhel8/support-tools:latest
-v /run:/run -v /var/log:/var/log -v /etc/machine-id:/etc/machine-id -v
/etc/localtime:/etc/localtime -v /:/host registry.redhat.io/rhel8/support-tools:latest
bash-4.4#
```

This opens a bash shell inside the support-tools container.

You can now run reports or debug tools against the host system (**/host**). When you are done, type **exit** to exit the shell and stop the container.

CHAPTER 6. BUILDING CONTAINER IMAGES WITH BUILDAH

The **buildah** command lets you create container images from a working container, a Dockerfile, or from scratch. The resulting images are OCI compliant, so they will work on any container runtime that meets the [OCI Runtime Specification](#) (such as Docker and CRI-O).

This section describes how to use the **buildah** command to create and otherwise work with containers and container images.

6.1. UNDERSTANDING BUILDAH

Using Buildah is different from building images with the **docker** command in the following ways:

- **No Daemon!:** Bypasses the Docker daemon! So no container runtime (Docker, CRI-O, or other) is needed to use Buildah.
- **Base image or scratch** Lets you not only build an image based on another container, but also lets you start with an empty image (scratch).
- **Build tools external:** Doesn't include build tools within the image itself. As a result, Buildah:
 - Reduces the size of images you build
 - Makes the image more secure by not having the software used to build the container (like gcc, make, and yum) within the resulting image.
 - Creates images that require fewer resources to transport the images (because they are smaller).

Buildah is able to operate without Docker or other container runtimes by storing data separately and by including features that let you not only build images, but run those images as containers as well. By default, Buildah stores images in an area identified as **containers-storage** (/var/lib/containers).



NOTE

The containers-storage location that the **buildah** command uses by default is the same place that the CRI-O container engine uses for storing local copies of images. So images pulled from a registry by either CRI-O or Buildah, or committed by the **buildah** command, will be stored in the same directory structure. Currently, however, CRI-O and Buildah cannot share containers, though they can share images.

There are more than a dozen options to use with the **buildah** command. Some of the main activities you can do with the **buildah** command include:

- **Build a container from a Dockerfile** Use a Dockerfile to build a new container image (**buildah bud**).
- **Build a container from another image or scratch** Build a new container, starting with an existing base image (**buildah from <imagename>**) or from scratch (**buildah from scratch**)
- **Inspecting a container or image** View metadata associated with the container or image (**buildah inspect**)
- **Mount a container:** Mount a container's root filesystem to add or change content (**buildah mount**).

- **Create a new container layer:** Use the updated contents of a container’s root filesystem as a filesystem layer to commit content to a new image (**buildah commit**).
- **Unmount a container:** Unmount a mounted container (**buildah umount**).
- **Delete a container or an image:** Remove a container (**buildah rm**) or a container image (**buildah rmi**).

For more details on Buildah, see the [GitHub Buildah page](#). The GitHub Buildah site includes man pages and software that might be more recent than is available with the RHEL version. Here are some other articles on Buildah that might interest you:

- [Buildah Tutorial 1: Building OCI container images](#)
- [Buildah Tutorial 2: Using Buildah with container registries](#)
- [Buildah Blocks - Getting Fit](#)

6.1.1. Installing Buildah

The buildah package is available with the container-tools module in RHEL 8 (**yum module install container-tools**). You can install the buildah package separately by typing:

```
# yum -y install buildah
```

With the buildah package installed, you can refer to the man pages included with the buildah package for details on how to use it. To see the available man pages and other documentation, then open a man page, type:

```
# rpm -qd buildah
# man buildah
buildah(1)      General Commands Manual      buildah(1)

NAME
Buildah - A command line tool that facilitates building OCI container images.
...
```

The following sections describe how to use **buildah** to get containers, build a container from a Dockerfile, build one from scratch, and manage containers in various ways.

6.2. GETTING IMAGES WITH BUILDDAH

To get a container image to use with **buildah**, use the **buildah from** command. Notice that if you are using RHEL 8.0, you may encounter problems with authenticating to the repository, see [bug](#). Here’s how to get a RHEL 8 image from the Red Hat Registry as a working container to use with the **buildah** command:

```
# buildah from registry.redhat.io/ubi8/ubi
Getting image source signatures
Copying blob...
Writing manifest to image destination
Storing signatures
ubi-working-container
# buildah images
```

```

IMAGE ID    IMAGE NAME                CREATED AT    SIZE
3da40a1670b5 registry.redhat.io/ubi8/ubi:latest May 8, 2019 21:55 214 MB
# buildah containers
CONTAINER ID BUILDER IMAGE ID    IMAGE NAME    CONTAINER NAME
c6c9279ecc0f *    3da40a1670b5 ...ubi8/ubi:latest ubi-working-container

```

Notice that the result of the **buildah from** command is an image (registry.redhat.io/ubi8/ubi:latest) and a working container that is ready to run from that image (ubi-working-container). Here's an example of how to execute a command from that container:

```

# buildah run ubi-working-container cat /etc/redhat-release
Red Hat Enterprise Linux release 8.0

```

The image and container are now ready for use with Buildah.

6.3. BUILDING AN IMAGE FROM A DOCKERFILE WITH BUILDDAH

With the **buildah** command, you can create a new image from a Dockerfile. The following steps show how to build an image that includes a simple script that is executed when the image is run.

This simple example starts with two files in the current directory: Dockerfile (which holds the instructions for building the container image) and myecho (a script that echoes a few words to the screen):

```

# ls
Dockerfile myecho
# cat Dockerfile
FROM registry.redhat.io/ubi8/ubi
ADD myecho /usr/local/bin
ENTRYPOINT "/usr/local/bin/myecho"
# cat myecho
echo "This container works!"
# chmod 755 myecho
# ./myecho
This container works!

```

With the Dockerfile in the current directory, build the new container as follows:

```

# buildah bud -t myecho .
STEP 1: FROM registry.redhat.io/ubi8/ubi
STEP 2: ADD myecho /usr/local/bin
STEP 3: ENTRYPOINT "/usr/local/bin/myecho"

```

The **buildah bud** command creates a new image named myecho. To see that new image, type:

```

# buildah images
IMAGE NAME    IMAGE TAG IMAGE ID    CREATED AT    SIZE
localhost/myecho latest    a3882af49784 Jun 21, 2019 12:21 216 MB

```

Next, you can run the image, to make sure it is working.

6.3.1. Running the image you built

To check that the image you built previously works, you can run the image using **podman run**:

-

```
# podman run localhost/myecho
This container works!
```

6.3.2. Inspecting a container with Buildah

With **buildah inspect**, you can show information about a container or image. For example, to build and inspect the **myecho** image, type:

```
# buildah from localhost/myecho
# buildah inspect localhost/myecho | less
{
  "Type": "buildah 0.0.1",
  "FromImage": "docker.io/library/myecho:latest",
  "FromImage-ID": "e2b190ac8...",
  "Config": "{\"created\":\"2018-11-13...

  "Entrypoint": [
    "/usr/local/bin/myecho"
  ],
  "WorkingDir": "/",
  "Labels": {
    "architecture": "x86_64",
    "authoritative-source-url": "registry.access.redhat.com",
    "build-date": "2018-09-19T20:46:28.459833",
```

To inspect a container from that same image, type the following:

```
# buildah inspect myecho-working-container | less
{
  "Type": "buildah 0.0.1",
  "FromImage": "docker.io/library/myecho:latest",
  "FromImage-ID": "e2b190a...",
  "Config": "{\"created\":\"2018-11-13T19:5...
...
  "Container": "myecho-working-container",
  "ContainerID": "c0cd2e494d...",
  "MountPoint": "",
  "ProcessLabel": "system_u:system_r:svirt_lxc_net_t:s0:c89,c921",
  "MountLabel": "",
```

Note that the container output has added information, such as the container name, container id, process label, and mount label to what was in the image.

6.4. MODIFYING A CONTAINER TO CREATE A NEW IMAGE WITH BUILDDAH

There are several ways you can modify an existing container with the **buildah** command and commit those changes to a new container image:

- Mount a container and copy files to it
- Use **buildah copy** and **buildah config** to modify a container

Once you have modified the container, use **buildah commit** to commit the changes to a new image.

6.4.1. Using buildah mount to modify a container

After getting an image with **buildah from**, you can use that image as the basis for a new image. The following text shows how to create a new image by mounting a working container, adding files to that container, then committing the changes to a new image.

Type the following to view the working container you used earlier:

```
# buildah containers
CONTAINER ID BUILDER IMAGE ID   IMAGE NAME   CONTAINER NAME

dc8f21af4a47 * 1456eedf8101 registry.redhat.io/ubi8/ubi:latest
                ubi-working-container
6d1ffccb557d * ab230ac5aba3 docker.io/library/myecho:latest
                myecho-working-container
```

Mount the container image and set the mount point to a variable (\$mymount) to make it easier to deal with:

```
# mymount=$(buildah mount myecho-working-container)
# echo $mymount
/var/lib/containers/storage/devicemapper/mnt/176c273fe28c23e5319805a2c48559305a57a706cc7ae7b
ec7da4cd79edd3c02/rootfs
```

Add content to the script created earlier in the mounted container:

```
# echo 'echo "We even modified it." >> $mymount/usr/local/bin/myecho
```

To commit the content you added to create a new image (named myecho), type the following:

```
# buildah commit myecho-working-container containers-storage:myecho2
```

To check that the new image includes your changes, create a working container and run it:

```
# buildah images
IMAGE ID   IMAGE NAME   CREATED AT      SIZE
a7e06d3cd0e2 docker.io/library/myecho2:latest
                Oct 12, 2017 15:15 3.144 KB
# buildah from docker.io/library/myecho2:latest
myecho2-working-container
# podman run docker.io/library/myecho2
This container works!
We even modified it.
```

You can see that the new **echo** command added to the script displays the additional text.

When you are done, you can unmount the container:

```
# buildah umount myecho-working-container
```

6.4.2. Using buildah copy and buildah config to modify a container

With **buildah copy**, you can copy files to a container without mounting it first. Here's an example, using the **myecho-working-container** created (and unmounted) in the previous section, to copy a new script to the container and change the container's configuration to run that script by default.

Create a script called **newecho** and make it executable:

```
# cat newecho
echo "I changed this container"
# chmod 755 newecho
```

Create a new working container:

```
# buildah from myecho:latest
myecho-working-container-2
```

Copy **newecho** to `/usr/local/bin` inside the container:

```
# buildah copy myecho-working-container-2 newecho /usr/local/bin
```

Change the configuration to use the **newecho** script as the new entrypoint:

```
# buildah config --entrypoint "/bin/sh -c /usr/local/bin/newecho" myecho-working-container-2
```

Run the new container, which should result in the **newecho** command being executed:

```
# buildah run myecho-working-container-2
I changed this container
```

If the container behaved as you expected it would, you could then commit it to a new image (mynewecho):

```
# buildah commit myecho-working-container-2 containers-storage:mynewecho
```

6.5. CREATING IMAGES FROM SCRATCH WITH BUILDAH

Instead of starting with a base image, you can create a new container that holds no content and only a small amount of container metadata. This is referred to as a **scratch** container. Here are a few issues to consider when choosing to create an image starting from a scratch container with the **buildah** command:

- When building a scratch container you can copy executable with no dependencies into the scratch image and make a few configuration settings to get a minimal container to work.
- To use tools like **yum** or **rpm** packages to populate the scratch container, you need to at least initialize an RPM database in the container and add a release package. The example below shows how to do that.
- If you end up adding a lot of RPM packages, consider using the **ubi** or **ubi-minimal** base images instead of a scratch image. Those base images have had documentation, language packs, and other components trimmed out, which can ultimately result in your image being smaller.

This example adds a Web service (httpd) to a container and configures it to run. To begin, create a scratch container:

```
# buildah from scratch
working-container
```

This creates just an empty container (no image) that you can mount as follows:

```
# scratchmnt=$(buildah mount working-container)
# echo $scratchmnt
/var/lib/containers/storage/devicemapper/mnt/cc92011e9a2b077d03a97c0809f1f3e7fef0f29bdc6ab5e86
b85430ec77b2bf6/rootfs
```

Initialize an RPM database within the scratch image and add the redhat-release package (which includes other files needed for RPMs to work):

```
# yum install -y --releasever=8 --installroot=$scratchmnt redhat-release
```

Install the httpd service to the scratch directory:

```
# yum install -y --setopt=reposdir=/etc/yum.repos.d \
--installroot=$scratchmnt \
--setopt=cachedir=/var/cache/dnf httpd
```

Add some text to an index.html file in the container, so you will be able to test it later:

```
# echo "Your httpd container from scratch worked." > $scratchmnt/var/www/html/index.html
```

Instead of running httpd as an init service, set a few **buildah config** options to run the httpd daemon directly from the container:

```
# buildah config --cmd "/usr/sbin/httpd -DFOREGROUND" working-container
# buildah config --port 80/tcp working-container
# buildah commit working-container localhost/myhttpd:latest
```

For now, you can use the Image ID to run the new image as a container with the **podman** command:

```
# podman images
REPOSITORY      TAG          IMAGE ID      CREATED      SIZE
localhost/myhttpd latest       47c0795d7b0e  9 minutes ago 665.6 MB
# podman run -p 8080:80 -d --name httpd-server 47c0795d7b0e
# curl localhost:8080
Your httpd container from scratch worked.
```

6.6. REMOVING IMAGES OR CONTAINERS WITH BUILDDAH

When you are done with particular containers or images, you can remove them with **buildah rm** or **buildah rmi**, respectively. Here are some examples.

To remove the container created in the previous section, you could type the following to see the mounted container, unmount it and remove it:

```
# buildah containers
CONTAINER ID  BUILDER  IMAGE ID  IMAGE NAME          CONTAINER NAME
05387e29ab93  *       c37e14066ac7 docker.io/library/myecho:latest myecho-working-container
```

```
# buildah mount
05387e29ab93 /var/lib/containers/storage/devicemapper/mnt/9274181773a.../rootfs
# buildah umount 05387e29ab93
# buildah rm 05387e29ab93
05387e29ab93151cf52e9c85c573f3e8ab64af1592b1ff9315db8a10a77d7c22
```

To remove the image you created previously, you could type the following:

```
# buildah rmi docker.io/library/myecho:latest
untagged: docker.io/library/myecho:latest
ab230ac5aba3b5a0a7c3d2c5e0793280c1a1b4d2457a75a01b70a4b7a9ed415a
```

6.7. USING CONTAINER REGISTRIES WITH BUILDDAH

With Buildah, you can push and pull container images between your local system and public or private container registries. The following examples show how to:

- Push containers to and pull them from a private registry with buildah.
- Push and pull container between your local system and the Docker Registry.
- Use credentials to associate your containers with a registry account when you push them.

Use the `skopeo` command, in tandem with the **buildah** command, to query registries for information about container images.

6.7.1. Pushing containers to a private registry

Pushing containers to a private container registry with the **buildah** command works much the same as pushing containers with the **docker** command. You need to:

- Set up a private registry (OpenShift provides a container registry or you can set up a Red Hat Quay container registry).
- Create or acquire the container image you want to push.
- Use **buildah push** to push the image to the registry.

To push an image from your local Buildah container storage, check the image name, then push it using the **buildah push** command. Remember to identify both the local image name and a new name that includes the location. For example, a registry running on the local system that is listening on TCP port 5000 would be identified as `localhost:5000`.

```
# buildah images
IMAGE ID    IMAGE NAME                CREATED AT    SIZE
cb702d492ee9 docker.io/library/myecho2:latest Nov 12, 2018 16:50    3.143 KB

# buildah push --tls-verify=false myecho2:latest localhost:5000/myecho2:latest
Getting image source signatures
Copying blob sha256:e4efd0...
...
Writing manifest to image destination
Storing signatures
```


Use the **curl** command to list the images in the registry and **skopeo** to inspect metadata about the image:

```
# curl http://localhost:5000/v2/_catalog
{"repositories":["myatomic","myecho2"]}
# curl http://localhost:5000/v2/myecho2/tags/list
{"name":"myecho2","tags":["latest"]}
# skopeo inspect --tls-verify=false docker://localhost:5000/myecho2:latest | less
{
  "Name": "localhost:5000/myecho2",
  "Digest": "sha256:8999ff6050...",
  "RepoTags": [
    "latest"
  ],
  "Created": "2017-11-21T16:50:25.830343Z",
  "DockerVersion": "",
  "Labels": {
    "architecture": "x86_64",
    "authoritative-source-url": "registry.redhat.io",
```

At this point, any tool that can pull container images from a container registry can get a copy of your pushed image. For example, on a RHEL 7 system you could start the docker daemon and try to pull the image so it can be used by the **docker** command as follows:

```
# systemctl start docker
# docker pull localhost:5000/myecho2
# docker run localhost:5000/myecho2
This container works!
```

6.7.2. Pushing containers to the Docker Hub

You can use your Docker Hub credentials to push and pull images from the Docker Hub with the **buildah** command. For this example, replace the username and password (testaccountXX:My00P@sswd) with your own Docker Hub credentials:

```
# buildah push --creds testaccountXX:My00P@sswd \
  docker.io/library/myecho2:latest docker://testaccountXX/myecho2:latest
```

As with the private registry, you can then get and run the container from the Docker Hub with the **podman**, **buildah** or **docker** command:

```
# podman run docker.io/textaccountXX/myecho2:latest
This container works!
# buildah from docker.io/textaccountXX/myecho2:latest
myecho2-working-container-2
# podman run myecho2-working-container-2
This container works!
```

CHAPTER 7. RUNNING CONTAINERS AS SYSTEMD SERVICES WITH PODMAN

Podman (Pod Manager) is a fully featured container engine that is a simple daemonless tool. Podman provides a Docker-CLI comparable command line that eases the transition from other container engines and allows the management of pods, containers and images. It was not originally designed to bring up an entire Linux system or manage services for such things as start-up order, dependency checking, and failed service recovery. That is the job of a full-blown initialization system like systemd.

Red Hat has become a leader in integrating containers with systemd, so that OCI and Docker-formatted containers built by Podman can be managed in the same way that other services and features are managed in a Linux system. This chapter describes how you can use the systemd initialization service to work with containers in two different ways:

- **Starting Containers with systemd:** By setting up a systemd unit file on your host computer, you can have the host automatically start, stop, check the status, and otherwise manage a container as a systemd service.
- **Starting services within a container using systemd** Many Linux services (Web servers, file servers, database servers, and so on) are already packaged for Red Hat Enterprise Linux to run as systemd services. If you are using the latest RHEL container image, you can set the RHEL container image to start the systemd service, then automatically start selected services within the container when the container starts up.

The following two sections describe how to use systemd container in those ways.

7.1. STARTING CONTAINERS WITH SYSTEMD

When you set up a container to start as a systemd service, you can define the order in which the containerized service runs, check for dependencies (like making sure another service is running, a file is available or a resource is mounted), and even have a container start by using the **runc** command.

This section provides an example of a container that is configured to run directly on a RHEL system as a systemd service. To learn about automatic generation of systemd service file, see [Generate systemd unit file](#).

1. Get the image you want to run on your system. For example, to use a minimal image based on Alpine Linux from docker.io, run the following command:

```
# podman pull docker.io/library/alpine:latest
```

2. Configure the container as a systemd service by creating the generic unit configuration file in the `~/.config/systemd/user` directory. For example, the contents of the `~/.config/systemd/user/container.service` can look as follows:

```
# cat ~/.config/systemd/user/container.service
[Unit]
Description=Podman in Systemd

[Service]
Restart=on-failure
ExecStartPre=/usr/bin/rm -f /%t/%n-pid /%t/%n-cid
ExecStart=/usr/bin/podman run --common-pidfile /%t/%n-pid --cidfile /%t/%n-cid -d
alpine:latest top
```

```
ExecStop=/usr/bin/sh -c "/usr/bin/podman rm -f `cat /%t/%n-cid`"
KillMode=none
Type=forking
PIDFile=/%t/%n-pid
```

```
[Install]
WantedBy=multi-user.target
```

- The **Restart=on-failure** line sets the restart policy and instructs systemd to restart the service when it cannot be started or stopped cleanly, or when the process exits with a non-zero status.
- The **ExecStart** line describes how we start the container.
- The **ExecStop** line describes how we stop and remove the container. You can run an **alpine:latest** container in the background that runs **top**. The **podman run** command includes two command-line options:
- The **--common-pidfile** option points to a path to store the process ID for the **common** process running on the host. The **common** process terminates with the same exit status as the container, which allows **systemd** to report the correct service status and restart the container if needed.
- The **--cidfile** option points to the path that stores the container ID.
- The **%t** is the path to the run time directory root, for example **/run/user/\$UserID**.
- The **%n** is the full name of the service.

For example, if the service name is **container** and the user ID is 1000, the above configuration places the **common-pidfile** in **/run/user/1000/container.service-pid** and the **cidfile** in **/run/user/1000/container.service-cid**.

3. To reload systemd manager configuration, type:

```
# systemctl --user daemon-reload
```

4. To enable the service and start it at boot time, type:

```
# systemctl --user enable container.service
```

5. To start the service immediately and check the status of the service, type the following:

```
# systemctl --user start container.service
# systemctl --user status container.service
* container.service - Podman in Systemd
   Loaded: loaded (/home/valentin/.config/systemd/user/container.service; disabled; vendor
   preset: enabled)
   Active: active (running) since Mon 2019-11-18 15:32:56 CET; 1min 5s ago
   Process: 189705 ExecStartPre=/usr/bin/rm -f //run/user/1000/container.service-pid
//run/user/1000/container.service-cid (code=exited, status=0/SUCCESS)
   Process: 189706 ExecStart=/usr/bin/podman run --common-pidfile
//run/user/1000/container.service-pid --cidfile //run/user/1000/container.service-cid -d
alpine:latest top (code=exited, status=0/SUCCESS)
   Main PID: 189731 (common)
   CGroup: /user.slice/user-1000.slice/user@1000.service/container.service
```

```

├─189724 /usr/bin/fuse-overlayfs [...]
├─189726 /usr/bin/slip4netns [...]
├─189731 /usr/bin/common [...]
└─189737 top

```

6. To list containers that are running or have exited, type:

```

# podman ps
CONTAINER ID  IMAGE                                COMMAND  CREATED   STATUS    PORTS
NAMES
f20988d59920  docker.io/library/alpine:latest  top     12 seconds ago  Up 11 seconds ago
funny_zhukovsky

```

7. To stop **container.service**, type:

```
# systemctl --user stop container.service
```

To learn more about configuring services with systemd, refer to the System Administrator's Guide chapter called [Managing Services with systemd](#) and article [Running containers with Podman and shareable systemd services](#).

7.2. STARTING SERVICES WITHIN A CONTAINER USING SYSTEMD

A package with the systemd initialization system is included in the official Red Hat Enterprise Linux Init base image named `registry.access.redhat.com/ubi8/ubi-init`. This means that applications created to be managed with systemd can be started and managed inside a container. A container running systemd will:

- Start the `/sbin/init` process (the systemd service) to run as PID 1 within the container.
- Start all systemd services that are installed and enabled within the container, in order of dependencies.
- Allow systemd to restart services or kill zombie processes for services started within the container.

The general steps for building a container that is ready to be used as a systemd services is:

- Install the package containing the systemd-enabled service inside the container. This can include dozens of services that come with RHEL, such as Apache Web Server (`httpd`), FTP server (`vsftpd`), Proxy server (`squid`), and many others. For this example, we simply install an Apache (`httpd`) Web server.
- The `httpd` and `vsftpd` packages are included in the UBI repositories. You would need a RHEL subscription to install the `squid` package.
- Use the `systemctl` command to enable the service inside the container.
- Add data for the service to use in the container (in this example, we add a Web server test page). For a real deployment, you would probably connect to outside storage.
- Expose any ports needed to access the service.

In this example, we build a container by creating a Dockerfile that installs and configures a Web server (`httpd`) to start automatically by the systemd service (`/sbin/init`) when the container is run on a host system.

1. **Create Dockerfile:** In a separate directory, create a file named Dockerfile with the following contents:

```
FROM registry.access.redhat.com/ubi8/ubi-init
RUN yum -y install httpd; yum clean all; systemctl enable httpd;
RUN echo "Successful Web Server Test" > /var/www/html/index.html
RUN mkdir /etc/systemd/system/httpd.service.d; echo -e '[Service]\nRestart=always' >
/etc/systemd/system/httpd.service.d/httpd.conf
EXPOSE 80
```

The Dockerfile installs the httpd package, enables the httpd service to start at boot time (i.e. when the container starts), creates a test file (index.html), exposes the Web server to the host (port 80), and starts the systemd init service (/sbin/init) when the container starts.

2. **Build the container:** From the directory containing the Dockerfile, type the following:

```
# podman build --format=docker -t mysysd .
```

3. **Open Selinux permission** If SELinux is enabled on your system, you must turn on the **container_manage_cgroup** boolean to run containers with systemd as shown here (see the [Containers running systemd](#) solution for details):

```
# setsebool -P container_manage_cgroup 1
```

4. **Run the container:** Once the container is built and named mysysd, type the following to run the container:

```
# podman run -d --name=mysysd_run -p 80:80 mysysd
```

From this command, the mysysd image runs as the mysysd_run container as a daemon process, with port 80 from the container exposed to port 80 on the host system.

5. **Check that the container is running** To make sure that the container is running and that the service is working, type the following commands:

```
# podman ps | grep mysysd_run
a282b0c2ad3d localhost/mysysd:latest /sbin/init 15 seconds ago Up 14 seconds ago
0.0.0.0:80->80/tcp mysysd_run
# curl localhost/index.html
Successful Web Server Test
```

At this point, you have a container that starts up a Web server as a systemd service inside the container. Install and run any services you like in this same way by modifying the Dockerfile and configuring data and opening ports as appropriate.

CHAPTER 8. CONTAINER COMMAND-LINE REFERENCE

8.1. PODMAN

The **podman** command (which stands for Pod Manager) lets you run containers as standalone entities, without requiring that Kubernetes, the Docker runtime, or any other container runtime be involved. It is a tool that can act as a replacement for the **docker** command, implementing the same command-line syntax, while it adds even more container management features. The **podman** features include:

- **Based on docker interface:** Because **podman** syntax mirrors the **docker** command, transitioning to **podman** should be easy for those familiar with **docker**.
- **Managing containers and images:** Both Docker- and OCI-compatible container images can be used with **podman** to:
 - Run, stop and restart containers
 - Create and manage container images (push, commit, configure, build, and so on)
- **Managing pods:** Besides running individual containers, **podman** can run a set of containers grouped in a pod. A pod is the smallest container unit that Kubernetes manages.
- **Working with no runtime:** No runtime environment is used by **podman** to work with containers.

Here are a few implementation features of **podman** you should know about:

- Podman, Buildah, and the CRI-O container engine all use the same back-end store directory, **/var/lib/containers**, instead of using the Docker storage location (**/var/lib/docker**), by default.
- Although Podman, Buildah, and CRI-O share the same storage directory, they cannot interact with each other's containers. Those tools can share images, however. Eventually those features will be able to share containers.
- The **podman** command, like the **docker** command, can build container images from a Dockerfile.
- The **podman** command can be a useful troubleshooting tool when the **CRI-O** service is unavailable.
- Options to the **docker** command that are not supported by **podman** include network, node, plugin (**podman** does not support plugins), rename (use **rm** and **create** to rename containers with **podman**), secret, service, stack, and swarm (**podman** does not support Docker Swarm). The container and image options are used to run subcommands that are used directly in **podman**.
- To interact programmatically with **podman**, [a remote API for Podman](#) is available using a technology called [varlink](#). This will let **podman** listen for API requests from remote tools (such as the RHEL 8 web console or the **atomic** command) and respond to them.

8.1.1. Using podman commands

If you are used to using the **docker** command to work with containers, you will find most of the features and options match those of **podman**. Table 1 shows a list of commands you can use with **podman** (type **podman -h** to see this list):

Table 8.1. Commands supported by podman

podman command	Description	podman command	Description
attach	Attach to a running container	commit	Create new image from changed container
build	Build an image using Dockerfile instructions	create	Create, but do not start, a container
diff	Inspect changes on container's filesystems	exec	Run a process in a running container
export	Export container's filesystem contents as a tar archive	help, h	Shows a list of commands or help for one command
history	Show history of a specified image	images	List images in local storage
import	Import a tarball to create a filesystem image	info	Display system information
inspect	Display the configuration of a container or image	kill	Send a specific signal to one or more running containers
load	Load an image from an archive	login	Login to a container registry
logout	Logout of a container registry	logs	Fetch the logs of a container
mount	Mount a working container's root filesystem	pause	Pauses all the processes in one or more containers
ps	List containers	port	List port mappings or a specific mapping for the container
pull	Pull an image from a registry	push	Push an image to a specified destination
restart	Restart one or more containers	rm	Remove one or more containers from host. Add -f if running.
rmi	removes one or more images from local storage	run	run a command in a new container

save	Save image to an archive	search	search registry for image
start	Start one or more containers	stats	Display percentage of CPU, memory, network I/O, block I/O and PIDs for one or more containers
stop	Stop one or more containers	tag	Add an additional name to a local image
top	Display the running processes of a container	umount, unmount	Unmount a working container's root filesystem
unpause	Unpause the processes in one or more containers	version	Display podman version information
wait	Block on one or more containers		

8.1.2. Trying basic podman commands

Because the use of **podman** mirrors the features and syntax of the **docker** command, you can refer to [Working with Docker Formatted Container Images](#) for examples of how to use those options to work with containers. Replace **docker** with **podman** in most cases. Here are some examples of using **podman**.

8.1.3. Pull a container image to the local system

```
# podman pull registry.redhat.io/ubi8/ubi
Trying to pull registry.redhat...Getting image source signatures
Copying blob sha256:d1fe25896eb5cbcee...
Writing manifest to image destination
Storing signatures
fd1ba0b398a82d56900bb798c...
```

8.1.4. List local container images

```
# podman images
REPOSITORY          TAG       IMAGE ID       CREATED        SIZE
registry.redhat.io/ubi8/ubi  latest   de9c26f23799  5 weeks ago   80.1MB
registry.redhat.io/ubi8/ubi  latest   fd1ba0b398a8  5 weeks ago   211MB
```

8.1.5. Inspect a container image

```
# podman inspect registry.redhat.io/ubi8/ubi | less
[
```



```
{
  "Id": "4bbd153adf8487a8a5114af0d6...",
  "Digest": "sha256:9999e735605c73f...",
  "RepoTags": [
    "registry.access.redhat.com/ubi8/ubi:latest"
  ],
  "RepoDigests": [
    "registry.access.redhat.com/ubi8/ubi/rhel@sha256:9999e7356..."
  ]
}
```

8.1.6. Run a container image

Run a container that opens a shell inside the container:

```
# podman run -it registry.redhat.io/ubi8/ubi /bin/bash
[root@8414218c04f9 /]# ps -ef
UID      PID  PPID  C  STIME TTY      TIME CMD
root      1    0  0 13:48 pts/0    00:00:00 /bin/bash
root     21    1  0 13:49 pts/0    00:00:00 ps -ef
[root@8414218c04f9 /]# exit
#
```

8.1.7. List containers that are running or have exited

```
# podman ps -a
CONTAINER ID  IMAGE                                COMMAND
CREATED AT      STATUS                                PORTS NAMES
440becd26893  registry.redhat.io/ubi8/ubi-minimal:latest /bin/bash
2018-05-10 09:02:52 -0400 EDT  Exited (0) About an hour ago  happy_hodgkin
8414218c04f9  registry.redhat.io/ubi8/ubi:latest      /bin/bash
2018-05-10 09:48:07 -0400 EDT  Exited (0) 14 minutes ago  nostalgic_boyd
```

8.1.8. Remove a container or image

Remove a container by its container ID:

```
# podman rm 440becd26893
```

Remove a container image by its image ID or name (use `-f` to force):

```
# podman rmi registry.redhat.io/ubi8/ubi
# podman rmi de9c26f23799
# podman rmi -f registry.redhat.io/ubi8/ubi:latest
```

8.1.9. Generate a Kube pod yaml file

Use the **podman generate** command to create a Kubernetes pod yaml file from a container or a pod:

1. Start a containerized service that runs as a daemon process (mariadb, in this example):

```
# podman run -d -e MYSQL_USER=user -e MYSQL_PASSWORD=pass \
-e MYSQL_DATABASE=db -p 3306:3306 --name mymariadb rhsc/mariadb-102-rhel7
```

2. Get the container name or ID:

```
# podman ps
CONTAINER ID  IMAGE
COMMAND      CREATED      STATUS
PORTS        NAMES
e421a3424ab0 registry.access.redhat.com/rhscv/mariadb-102-rhel7:latest
container-entrypo... 19 seconds ago Up 16 seconds ago
0.0.0.0:3306->3306/tcp mymariadb
```

3. Use the container name or ID to generate the Kubernetes yaml output and direct it into a file:

```
# podman generate kube mymariadb > mymariadbkube.yaml
# less mymariadbkube.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2019-10-31T13:19:41Z"
  labels:
    app: mymariadb
    name: mymariadb
spec:
  containers:
    - command:
      ...
    - name: MYSQL_VERSION
      value: "10.2"
    - name: CONTAINER_SCRIPTS_PATH
      value: /usr/share/container-scripts/mysql
    - name: MYSQL_PREFIX
      value: /opt/rh/rh-mariadb102/root/usr
    - name: ENABLED_COLLECTIONS
      value: rh-mariadb102
    - name: DESCRIPTION
      value: MariaDB is a multi-user, multi-threaded SQL database server. The container
        image provides a containerized packaging of the MariaDB mysqld daemon and
        client application. The mysqld server daemon accepts connections from clients
        and provides access to content from MariaDB databases on behalf of the clients.
    - name: STI_SCRIPTS_PATH
      value: /usr/libexec/s2i
    - name: PROMPT_COMMAND
      value: . /usr/share/container-scripts/mysql/scl_enable
    - name: MYSQL_USER
      value: user
    - name: MYSQL_PASSWORD
      value: pass
  ...
```

The **podman generate** command does not reflect any Logical Volume Manager (LVM) logical volumes or physical volumes that might be attached to the container.

4. You can now use the yaml file to create a pod in your Kubernetes or OpenShift environment:

```
# kubectl create -f mymariadbkube.yaml
```

8.1.10. Generate a systemd unit file

Podman allows **systemd** to control and manage container processes. You can generate a systemd unit file for an existing container as non-root (rootless) user using the **podman generate systemd** command:

1. Create the container (e.g. **myfedora**):

```
# podman create -d --name myfedora fedora:latest top
54502f309f3092d32b4c496ef3d099b270b2af7b5464e7cb4887bc16a4d38597
```

2. Use the container name or ID to generate the systemd unit file and direct it into the `~/.config/systemd/user/container-myfedora.service` file:

```
# podman generate systemd --name myfedora > ~/.config/systemd/user/container-
myfedora.service
container-myfedora.service
autogenerated by Podman 1.6.2
Tue Nov 19 15:49:15 CET 2019

[Unit]
Description=Podman container-myfedora.service
Documentation=man:podman-generate-systemd(1)

[Service]
Restart=on-failure
ExecStart=/usr/bin/podman start myfedora
ExecStop=/usr/bin/podman stop -t 10 myfedora
KillMode=none
Type=forking
PIDFile=/run/user/1000/overlay-
containers/54502f309f3092d32b4c496ef3d099b270b2af7b5464e7cb4887bc16a4d38597/userd
ata/common.pid

[Install]
WantedBy=multi-user.target
```

3. You can now start and check the status of the container:

```
# systemctl --user start container-myfedora.service
# systemctl --user status container-myfedora.service
```

4. You can stop the container:

```
# systemctl --user stop container-myfedora.service
```

For more information, see article [Running containers with Podman and shareable systemd services](#) .

8.1.11. Creating SELinux policies for containers

To generate SELinux policies for containers, use the UDICA tool. For more information, see [Introduction to the udica SELinux policy generator](#).

8.1.12. Using podman with MPI

You can use Podman with Open MPI (Message Passing Interface) to run containers in a High Performance Computing (HPC) environment.

The example is based on the [ring.c](#) program taken from Open MPI. In this example, a value is passed around by all processes in a ring-like fashion. Each time the message passes rank 0, the value is decremented. When each process receives the 0 message, it passes it on to the next process and then quits. By passing the 0 first, every process gets the 0 message and can quit normally.

Procedure

1. Install Open MPI:

```
$ sudo yum install openmpi
```

2. To activate the environment modules, type:

```
$ . /etc/profile.d/modules.sh
```

3. Load the **mpi/openmpi-x86_64** module:

```
$ module load mpi/openmpi-x86_64
```

Optionally, to automatically load **mpi/openmpi-x86_64** module, add this line to the **.bashrc** file:

```
$ echo "module load mpi/openmpi-x86_64" >> .bashrc
```

4. To combine **mpirun** and **podman**, create a container with the following definition:

```
$ cat Containerfile
FROM registry.access.redhat.com/ubi8/ubi

RUN yum -y install openmpi-devel wget && \
    yum clean all

RUN wget https://raw.githubusercontent.com/open-mpi/ompi/master/test/simple/ring.c && \
    /usr/lib64/openmpi/bin/mpicc ring.c -o /home/ring && \
    rm -f ring.c
```

5. Build the container:

```
$ podman build --tag=mpi-ring .
```

6. Start the container. On a system with 4 CPUs this command starts 4 containers:

```
$ mpirun \
  --mca orte_tmpdir_base /tmp/podman-mpirun \
  podman run --env-host \
    -v /tmp/podman-mpirun:/tmp/podman-mpirun \
    --userns=keep-id \
    --net=host --pid=host --ipc=host \
    mpi-ring /home/ring
Rank 2 has cleared MPI_Init
Rank 2 has completed ring
```

```

Rank 2 has completed MPI_Barrier
Rank 3 has cleared MPI_Init
Rank 3 has completed ring
Rank 3 has completed MPI_Barrier
Rank 1 has cleared MPI_Init
Rank 1 has completed ring
Rank 1 has completed MPI_Barrier
Rank 0 has cleared MPI_Init
Rank 0 has completed ring
Rank 0 has completed MPI_Barrier

```

As a result, **mpirun** starts up 4 Podman containers and each container is running one instance of the **ring** binary. All 4 processes are communicating over MPI with each other.

The following **mpirun** options are used to start the container:

- **--mca orte_tmpdir_base /tmp/podman-mpirun** line tells Open MPI to create all its temporary files in **/tmp/podman-mpirun** and not in **/tmp**. If using more than one node this directory will be named differently on other nodes. This requires mounting the complete **/tmp** directory into the container which is more complicated.

The **mpirun** command specifies the command to start, the **podman** command. The following **podman** options are used to start the container:

- **run** command runs a container.
- **--env-host** option copies all environment variables from the host into the container.
- **-v /tmp/podman-mpirun:/tmp/podman-mpirun** line tells Podman to mount the directory where Open MPI creates its temporary directories and files to be available in the container.
- **--userns=keep-id** line ensures the user ID mapping inside and outside the container.
- **--net=host --pid=host --ipc=host** line sets the same network, PID and IPC namespaces.
- **mpi-ring** is the name of the container.
- **/home/ring** is the MPI program in the container.

For more information, see the article [Podman in HPC environments](#) by Adrian Reber.

8.1.13. Creating and restoring container checkpoints

Checkpoint/Restore In Userspace (CRIU) is a software that enables you to set a checkpoint on a running container or an individual application and store its state to disk. You can use data saved to restore the container after a reboot at the same point in time it was checkpointed.

8.1.13.1. Creating and restoring a container checkpoint locally

This example is based on a Python based web server which returns a single integer which is incremented after each request.

Procedure

1. Create a Python based server:

```
# cat counter.py
#!/usr/bin/python3

import http.server

counter = 0

class handler(http.server.BaseHTTPRequestHandler):
    def do_GET(s):
        global counter
        s.send_response(200)
        s.send_header('Content-type', 'text/html')
        s.end_headers()
        s.wfile.write(b'%d\n' % counter)
        counter += 1

server = http.server.HTTPServer(("", 8088), handler)
server.serve_forever()
```

2. Create a container with the following definition:

```
# cat Containerfile
FROM registry.access.redhat.com/ubi8/ubi

COPY counter.py /home/counter.py

RUN useradd -ms /bin/bash counter

RUN yum -y install python3 && chmod 755 /home/counter.py

USER counter
ENTRYPOINT /home/counter.py
```

The container is based on the Universal Base Image (UBI 8) and uses a Python based server.

3. Build the container:

```
# podman build . --tag counter
```

Files **counter.py** and **Containerfile** are the input for the container build process (**podman build**). The built image is stored locally and tagged with the tag **counter**.

4. Start the container as root:

```
# podman run --name criu-test --detach counter
```

5. To list all running containers, enter:

```
# podman ps
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES
e4f82fd84d48 localhost/counter:latest 5 seconds ago Up 4 seconds ago criu-test
```

6. Display IP address of the container:

```
# podman inspect criu-test --format "{{.NetworkSettings.IPAddress}}"
10.88.0.247
```

7. Send requests to the container:

```
# curl 10.88.0.247:8080
0
# curl 10.88.0.247:8080
1
```

8. Create a checkpoint for the container:

```
# podman container checkpoint criu-test
```

9. Reboot the system.

10. Restore the container:

```
# podman container restore --keep criu-test
```

11. Send requests to the container:

```
# curl 10.88.0.247:8080
2
# curl 10.88.0.247:8080
3
# curl 10.88.0.247:8080
4
```

The result now does not start at **0** again, but continues at the previous value.

This way you can easily save the complete container state through a reboot.

For more information, see the article [Adding checkpoint/restore support to Podman](#) by Adrian Reber.

8.1.13.2. Reducing startup time using container restore

You can use container migration to reduce startup time of containers which require a certain time to initialize. Using a checkpoint, you can restore the container multiple times on the same host or on different hosts. This example is based on the container from the [Creating and restoring a container checkpoint locally](#) section.

Procedure

1. Create a checkpoint of the container, and export the checkpoint image to a **tar.gz** file:

```
# podman container checkpoint criu-test --export /tmp/chkpt.tar.gz
```

2. Restore the container from the **tar.gz** file:

```
# podman container restore --import /tmp/chkpt.tar.gz --name counter1
# podman container restore --import /tmp/chkpt.tar.gz --name counter2
# podman container restore --import /tmp/chkpt.tar.gz --name counter3
```

The **--name (-n)** option specifies a new name for containers restored from the exported checkpoint.

3. Display ID and name of each container:

```
# podman ps -a --format "{{.ID}} {{.Names}}"
a8b2e50d463c counter3
faabc5c27362 counter2
2ce648af11e5 counter1
```

4. Display IP address of each container:

```
# podman inspect counter1 --format "{{.NetworkSettings.IPAddress}}"
10.88.0.248

# podman inspect counter2 --format "{{.NetworkSettings.IPAddress}}"
10.88.0.249

# podman inspect counter3 --format "{{.NetworkSettings.IPAddress}}"
10.88.0.250
```

5. Send requests to each container:

```
# curl 10.88.0.248:8080
4
# curl 10.88.0.249:8080
4
# curl 10.88.0.250:8080
4
```

Note, that the result is **4** in all cases, because you are working with different containers restored from the same checkpoint.

Using this approach, you can quickly start up stateful replicas of the initially checkpointed container.

For more information, see the article [Container migration with Podman on RHEL](#) by Adrian Reber.

8.1.13.3. Migrating containers among systems

This procedure shows the migration of running containers from one system to another, without losing the state of the applications running in the container. This example is based on the container from the [Creating and restoring a container checkpoint locally](#) section tagged with **counter**.

Prerequisites

The following steps are not necessary if the container is pushed to a registry as Podman will automatically download the container from a registry if it is not available locally. This example does not use a registry, you have to export previously built and tagged container (see [Creating and restoring a container checkpoint locally](#) section) locally and import the container on the destination system of this migration.

- Export previously built container:

```
# podman save --output counter.tar counter
```


- Copy exported container image to the destination system (***other_host***):

```
# scp counter.tar other_host:
```

- Import exported container on the destination system:

```
# ssh other_host podman load --input counter.tar
```

Now the destination system of this container migration has the same container image stored in its local container storage.

Procedure

1. Start the container as root:

```
# podman run --name criu-test --detach counter
```

2. Display IP address of the container:

```
# podman inspect criu-test --format "{{.NetworkSettings.IPAddress}}"
10.88.0.247
```

3. Send requests to the container:

```
# curl 10.88.0.247:8080
0
# curl 10.88.0.247:8080
1
```

4. Create a checkpoint of the container, and export the checkpoint image to a **tar.gz** file:

```
# podman container checkpoint criu-test --export /tmp/chkpt.tar.gz
```

5. Copy the checkpoint archive to the destination host:

```
# scp /tmp/chkpt.tar.gz other_host:/tmp/
```

6. Restore the checkpoint on the destination host (***other_host***):

```
# podman container restore --import /tmp/chkpt.tar.gz
```

7. Send a request to the container on the destination host (***other_host***):

```
# curl 10.88.0.247:8080
2
```

As a result, the stateful container has been migrated from one system to another without losing its state.

For more information, see the article [Container migration with Podman on RHEL](#) by Adrian Reber.

8.2. RUNC

"runC" is a lightweight, portable implementation of the Open Container Initiative (OCI) container runtime specification. runC unites a lot of the low-level features that make running containers possible. It shares a lot of low-level code with Docker but it is not dependent on any of the components of the Docker platform.

runc supports Linux namespaces, live migration, and has portable performance profiles. It also provides full support for Linux security features such as SELinux, control groups (cgroups), seccomp, and others. You can build and run images with runc, or you can run OCI-compatible images with runc.

8.2.1. Running containers with runc

With runc, containers are configured using bundles. A bundle for a container is a directory that includes a specification file named "config.json" and a root filesystem. The root filesystem contains the contents of the container.

To create a bundle, run:

```
$ runc spec
```

This command creates a config.json file that only contains a bare-bones structure that you will need to edit. Most importantly, you will need to change the "args" parameter to identify the executable to run. By default, "args" is set to "sh".

```
"args": [  
  "sh"  
],
```

As an example, you can download the Red Hat Enterprise Linux base image (ubi8/ubi) using podman then export it, create a new bundle for it with runc, and edit the "config.json" file to point to that image. You can then create the container image and run an instance of that image with runc. Use the following commands:

```
# podman pull registry.redhat.io/ubi8/ubi  
# podman export $(podman create registry.redhat.io/ubi8/ubi) > rhel.tar  
# mkdir -p rhel-runc/rootfs  
# tar -C rhel-runc/rootfs -xf rhel.tar  
# runc spec -b rhel-runc  
# vi rhel-runc/config.json  Change any setting you like  
# runc create -b rhel-runc/ rhel-container  
# runc start rhel-container  
sh-4.2#
```

In this example, the name of the container instance is "rhel-container". Running that container, by default, starts a shell, so you can begin looking around and running commands from inside that container. Type **exit** when you are done.

The name of a container instance must be unique on the host. To start a new instance of a container:

```
# runc start <container_name>
```

You can provide the bundle directory using the "-b" option. By default, the value for the bundle is the current directory.

You will need root privileges to start containers with `runc`. To see all commands available to `runc` and their usage, run `"runc --help"`.

8.3. SKOPEO

With the `skopeo` command, you can work with container images from registries without using the `docker` daemon or the **`docker`** command. Registries can include the Docker Registry, your own local registries, Red Hat Quay or OpenShift registries. Activities you can do with `skopeo` include:

- **inspect:** The output of a **`skopeo inspect`** command is similar to what you see from a **`docker inspect`** command: low-level information about the container image. That output can be in json format (default) or raw format (using the `--raw` option).
- **copy:** With **`skopeo copy`** you can copy a container image from a registry to another registry or to a local directory.
- **layers:** The **`skopeo layers`** command lets you download the layers associated with images so that they are stored as tarballs and associated manifest files in a local directory.


Like the **`buildah`** command and other tools that rely on the `containers/image` library, the **`skopeo`** command can work with images from container storage areas other than those associated with Docker. Available transports to other types of container storage include: `containers-storage` (for images stored by **`buildah`** and CRI-O), `ostree` (for atomic and system containers), `oci` (for content stored in an OCI-compliant directory), and others. See the [skopeo man page](#) for details.

To try out `skopeo`, you could set up a local registry, then run the commands that follow to inspect, copy, and download image layers. If you want to follow along with the examples, start by doing the following:

- Install a local registry (such as [Red Hat Quay](#). Container registry software available in the `docker-distribution` package for RHEL 7, is not available for RHEL 8).
- Pull the latest RHEL image to your local system (**`podman pull ubi8/ubi`**).
- Retag the RHEL image and push it to your local registry as follows:

```
# podman tag ubi8/ubi localhost/myubi8
# podman push localhost/myubi8
```

The rest of this section describes how to inspect, copy and get layers from the RHEL image.



NOTE

The **`skopeo`** tool by default requires a TLS connection. It fails when trying to use an unencrypted connection. To override the default and use an `http` registry, prepend **`http:`** to the `<registry>/<image>` string.

8.3.1. Inspecting container images with skopeo

When you inspect a container image from a registry, you need to identify the container format (such as `docker`), the location of the registry (such as `docker.io` or `localhost`), and the repository/image (such as `ubi8/ubi`).

The following example inspects the `mariadb` container image from the Docker Registry:

```
# skopeo inspect docker://docker.io/library/mariadb
```

```
{
  "Name": "docker.io/library/mariadb",
  "Tag": "latest",
  "Digest": "sha256:d3f56b143b62690b400ef42e876e628eb5e488d2d0d2a35d6438a4aa841d89c4",
  "RepoTags": [
    "10.0.15",
    "10.0.16",
    "10.0.17",
    "10.0.19",
  ],
  ...
  "Created": "2018-06-10T01:53:48.812217692Z",
  "DockerVersion": "1.10.3",
  "Labels": {},
  "Architecture": "amd64",
  "Os": "linux",
  "Layers": [
  ],
  ...
}
```

Assuming you pushed a container image tagged **localhost/myubi8** to a container registry running on your local system, the following command inspects that image:

```
# skopeo inspect docker://localhost/myubi8
{
  "Name": "localhost/myubi8",
  "Tag": "latest",
  "Digest": "sha256:4e09c308a9ddf56c0ff6e321d135136eb04152456f73786a16166ce7cba7c904",
  "RepoTags": [
    "latest"
  ],
  "Created": "2018-06-16T17:27:13Z",
  "DockerVersion": "1.7.0",
  "Labels": {
    "Architecture": "x86_64",
    "Authoritative_Registry": "registry.access.redhat.com",
    "BZComponent": "rhel-server-docker",
    "Build_Host": "rcm-img01.build.eng.bos.redhat.com",
    "Name": "myubi8",
    "Release": "75",
    "Vendor": "Red Hat, Inc.",
    "Version": "8.0"
  },
  "Architecture": "amd64",
  "Os": "linux",
  "Layers": [
    "sha256:16dc1f96e3a1bb628be2e00518fec2bb97bd5933859de592a00e2eb7774b6ecf"
  ]
}
```

8.3.2. Copying container images with skopeo

This command copies the myubi8 container image from a local registry into a directory on the local system:

```
# skopeo copy docker://localhost/myubi8 dir:/root/test/
INFO[0000] Downloading
```

```
myubi8/blobs/sha256:16dc1f96e3a1bb628be2e00518fec2bb97bd5933859de592a00e2eb7774b6ecf
# ls /root/test
16dc1f96e3a1bb628be2e00518fec2bb97bd5933859de592a00e2eb7774b6ecf.tar manifest.json
```

The result of the **skopeo copy** command is a tarball (16d*.tar) and a manifest.json file representing the image being copied to the directory you identified. If there were multiple layers, there would be multiple tarballs. The **skopeo copy** command can also copy images to another registry. If you need to provide a signature to write to the destination registry, you can do that by adding a **--sign-by=** option to the command line, followed by the required key-id.

8.3.3. Getting image layers with skopeo

The **skopeo layers** command is similar to **skopeo copy**, with the difference being that the **copy** option can copy an image to another registry or to a local directory, while the **layers** option just drops the layers (tarballs and manifest.json file) in the current directory. For example

```
# skopeo layers docker://localhost/myubi8
INFO[0000] Downloading
myubi8/blobs/sha256:16dc1f96e3a1bb628be2e00518fec2bb97bd5933859de592a00e2eb7774b6ecf
# find .
./layers-myubi8-latest-698503105
./layers-myubi8-latest-698503105/manifest.json
./layers-myubi8-latest-
698503105/16dc1f96e3a1bb628be2e00518fec2bb97bd5933859de592a00e2eb7774b6ecf.tar
```

As you can see from this example, a new directory is created (layers-myubi8-latest-698503105) and, in this case, a single layer tarball and a manifest.json file are copied to that directory.

CHAPTER 9. ADDITIONAL RESOURCES

- [Buildah](#) - a tool for building OCI container images
- [Podman](#) - a tool for running and managing containers
- [Skopeo](#) - a tool for copying and inspecting container images